OPAL: AN OBSERVED-SINGLE BENCHMARK FOR COMBINATORIAL PERTURBATION PREDICTION

RESEARCH PREPRINT

Arjun Vijay Prakash
Independent researcher
High school student, City Montessori School

ABSTRACT

Predicting an unmeasured intervention pair from matched measured singles is a narrow population-mean task, and improvement over additivity need not establish pair-specific learning. Opal provides symmetric residual predictors and controlled comparisons with strong same-information references. Historical neural models reduce mean squared error over additivity by 15.81% on Norman and 61.08% on Wessels, but their advantage over development-selected simple references is inconclusive or negative. Exposed development favors economical residual kernels, and increased pair coverage helps more than the tested increase in single coverage on a fixed Wessels validation population. Technical partitions reveal unstable residual estimates without establishing a biological accuracy ceiling. An independent D-SPIN procedure evaluation yields MSE 0.034025 versus 0.033108 for the source-selected reference, failing superiority; source uncertainty also fails under shift. Preserved continuations support narrow known-compound dose interpolation but no newly validated pair/state component. A subsequent finite model-series evaluation trains main-effect selection, symmetric cross-coordinate bilinear residuals, and nested residual-memory routing. Bilinear improvements over the stronger Base reference are below 0.1% on both responsive primary panels; routing fails its promotion rule. Three upstream-method adaptations execute under disclosed input and protocol differences, without establishing native competitiveness. The release supplies reproducible comparative evidence and scoped simple references, while positive independent transfer, biological interaction recovery, and calibrated uncertainty remain unestablished.

Keywords: combinatorial perturbations; transcriptomic prediction; residual learning; symmetric models; reproducibility; evaluation.

1 Introduction

Single-cell perturbation experiments connect interventions to transcriptional responses, while combinatorial screens measure departures from single-intervention behavior. Norman's K562 CRISPR activation experiment provides a widely used genetic combination dataset (Norman et al., 2019). CaRPool-seq extends combinatorial targeting to RNA transcripts using Cas13 and provides a second experimental system (Wessels et al., 2023). These experiments motivate a prediction task that is narrower than general perturbation prediction: given measured single profiles in one context, predict an unmeasured simultaneous pair.

The additive prediction supplies a strong starting point. A learned model can improve on it by capturing a shared calibration error, a common offset, or pair-specific structure. Only the last explanation supports a claim that particular pair features contribute useful information. Metric calibration also matters because errors across many weakly responding genes and correlations dominated by single effects can obscure the quality of interaction prediction (Ahlmann-Eltze et al., 2025; Miller et al., 2026). We therefore compare the proposed correction to stronger controls using the same observed singles, and assess synthetic negative controls before interpreting biological gains.

This work implements Opal as a family of output-level residual models and an auditable experimental pipeline. The present contribution is the controlled internal comparison and its negative biological finding, rather than a demonstrated state-of-the-art method. Residual subtraction, symmetry, and dimensionality restriction are established modeling tools.

The experiments ask which of these restrictions help under a frozen selection procedure, and retain unfavorable results. The complete research scope also requires independent genetic confirmation, chemical validation, public-method comparisons, and fuller uncertainty assessment; the current manuscript records their incomplete status.

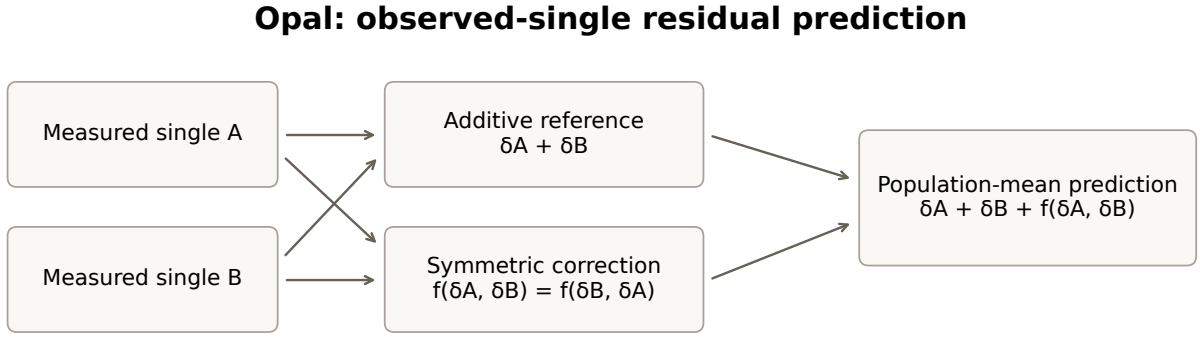


Figure 1: The observed-single prediction contract. A symmetric correction is added to two measured effects; its residual target includes shared distortion and measurement error as well as possible pair-specific structure.

2 Related work

CPA learns compositional representations for predicting responses to complex perturbations and supplies the earlier chemical feasibility source considered here (Lotfollahi et al., 2023). GEARS uses biological knowledge graphs to predict multigene perturbation responses (Roohani et al., 2024). Neither a published score under a different split nor access to external knowledge is interchangeable with a comparison that supplies observed singles under our protocol. Our internal controls use the same inputs. Later native CPA and GEARS aggregate adaptations are reported separately below, with their departures from canonical cell-level evaluation disclosed.

NAIAD combines single-effect information and learned representations for combinatorial experimental discovery (Qin et al., 2025). Its phenotype and acquisition objectives constrain claims of novelty for any single-informed pair model, but a transcriptomic comparison requires compatible outputs and data access. SLIM supplies a small linear reference using STRING embeddings (Hu et al., 2026); its upstream resource must be accounted for separately. PertAdapt adapts single-cell foundation models to perturbation prediction (Bai et al., 2026), so upstream pretraining and combination-task compatibility must also be disclosed. These references motivate future controlled benchmarks; no result in this manuscript ranks Opal against these methods.

3 Methods

3.1 Task and expression scale

Let $\mu_{u,s,g}$ denote the mean library-normalized count for intervention u, matched stratum s, and gene g, and let $\mu_{0,s,g}$ denote its control mean. Each cell is normalized to the median retained control library total. The primary effect is

$$\delta_{u,g} = \frac{1}{|S_u|} \sum_{s \in S_u} \left[\log_2(1 + \mu_{u,s,g}) - \log_2(1 + \mu_{0,s,g}) \right]. \tag{1}$$

We transform condition means, rather than averaging transformed cell counts. Strata contribute equally, and a retained pair must share stratum support with both component singles. The source-specific stratum definitions are technical matching variables, not assumed independent biological replicates.

For two distinct simultaneous interventions A and B, define $a_{AB} = \delta_A + \delta_B$ and $r_{AB} = \delta_{AB} - a_{AB}$. The prediction is

$$\widehat{\delta}_{AB} = a_{AB} + f_\theta(\delta_A, \delta_B), \qquad f_\theta(\delta_A, \delta_B) = f_\theta(\delta_B, \delta_A). \tag{2}$$

In the historical genetic benchmark, all components have observed single profiles in the same experimental context. That protocol does not evaluate unseen singles, different cell types, different doses, sequential treatments, or individual-

cell trajectories. Additivity depends on the chosen expression scale, and r_{AB} also includes error in the estimated singles and controls. It is a prediction residual, not automatically drug synergy or a causal genetic mechanism.

Residual and response losses are identical when the observed additive term is fixed:

$$\|f_\theta - r_{AB}\|_2^2 = \|a_{AB} + f_\theta - \delta_{AB}\|_2^2. \tag{3}$$

Any benefit from residual training therefore arises from parameterization, initialization, or regularization, rather than a different squared-error objective.

3.2 Source preparation and input access

Norman uses the named secondary scPerturb release from Zenodo record 7278143 version 1.1 (Peidli et al., 2024). Its documented transformation reads the original count matrix without expression normalization. We retain source-annotated good-coverage singlet assignments and explicitly negative-control constructs, preserve guide provenance, and match by gemgroup. This is distinct from a GEARS compatibility track. The initial library range is 500–20,000 counts, with at least 50 cells per condition. After source-specific exclusions and target-pair aggregation, the prepared dataset has 105 singles, 126 pairs, and 10,623 eligible genes; 88 pairs are development conditions and 38 are final test conditions.

Wessels uses the original GSE213957 four-lane archive and THP1 metadata (Wessels et al., 2023). Barcode joins retain metadata-matched THP1 Gene Expression measurements and exclude HEK293 cells and guide-count features. Exact `NT` tokens determine controls, avoiding substring misclassification of target names. The minimum library is 500; the maximum is the 99.5th percentile of controls and singles. Conditions require at least 50 cells, and lanes supply matched strata. Many measured singles fail this minimum, so the accepted task covers 16 singles and 58 pairs, with 10,590 genes, 40 development pairs, and 18 test pairs. This restricted coverage is a consequence of the observed-single contract.

Gene eligibility uses controls and singles only. A gene survives if at least one reference condition has normalized mean at least 0.1 and detection fraction at least 0.01. This frozen universe is the primary output; it does not establish performance on genes excluded by the reference rule. Source hashes, count locations, guide mapping, exclusions, normalization parameters, gene order, and split manifests are retained. Held-out pair responses occupy a separate evaluator bundle and are unavailable to model fitting or architecture selection.

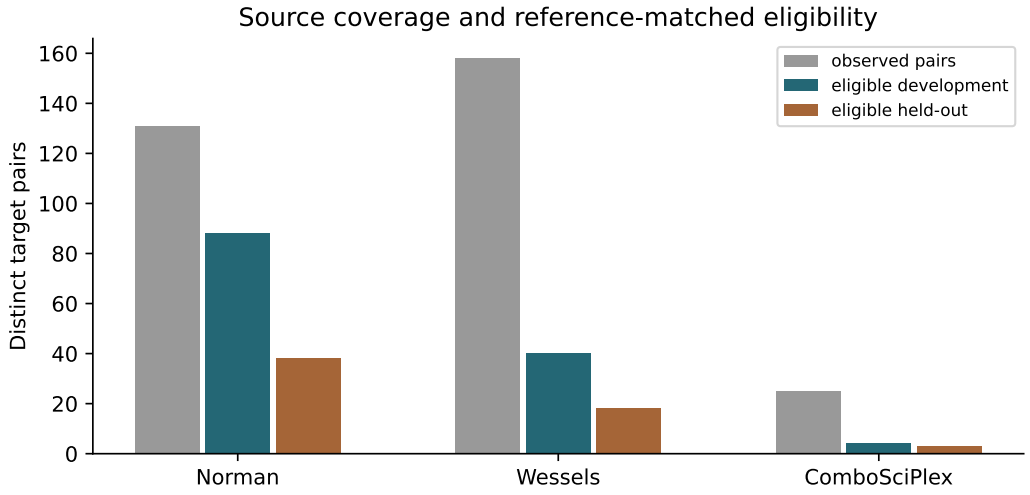


Figure 2: Original distinct nonrepeat pair coverage and surviving reference-matched development/test conditions. Source-level measurement does not guarantee a usable pair with both matched singles.

3.3 Representations and correction families

PCA fitted to centered single effects supplies an embedding $e_u \in \mathbb{R}^m$. Its dimension is capped by the number of available singles and output genes; nominal dimensions are 32 for Norman and 16 for Wessels and the synthetic suite. An identity representation and a linear denoising encoder are comparison candidates. When split-half single profiles

exist, the encoder learns reciprocal reconstruction between halves. Otherwise, the synthetic comparison uses added noise. Joint fine-tuning updates the target embedding table with a single-profile reconstruction penalty; it is not a demonstrated encoder for previously unmeasured targets.

The coordinate-product model uses $q_{AB} = e_A \odot e_B$. A learned bilinear candidate uses

$$q_{AB} = (Pe_A) \odot (Pe_B), \qquad f_\theta = b + B\Lambda^\top q_{AB}, \tag{4}$$

where $B \in \mathbb{R}^{G\times k}$ defines a shared output space. An independent symmetric-form candidate uses coordinates $e_A^\top W_j e_B$ with symmetric W_j. The neural network candidate applies a tanh network to concatenated $e_A + e_B$, $|e_A - e_B|$, and projected products. Its sum-network control omits the projected-product block but retains the sum and absolute difference. We therefore call this control a network without multiplicative features, rather than claiming that it receives only the sum. Both networks predict coefficients in the output basis plus an optional offset. Unrestricted-output, fixed-basis, zero-offset, random-initialization, direct-response, representation, and fine-tuning comparisons test individual choices.

3.4 Baselines and development selection

The internal baseline set includes no-change, additive, training-mean response, training-mean residual, symmetric ridge regression, reduced-rank regression, nearest-neighbor residual correction, and global correction. Symmetric regression uses sum, absolute difference, and coordinate products of the same embeddings. Global correction uses the observed additive response directly, with an optional feature expansion comprising its coordinates, their hyperbolic tangents, and their squares. It can explain shared nonlinear distortions without using pair identities.

Random pair holdout reserves 30% of canonical unordered pairs, using identities alone. Five development folds determine candidate errors and stopping epochs. The initial CPU search uses ranks 2 and 8, ridge penalties 1 and 100, and at most 80 epochs for neural candidates. AdamW uses learning rate 0.01 and weight decay 0.01, with full-batch equal-condition loss. Initial screening uses one optimization seed; real-data finalist evaluation uses five seeds, while synthetic engineering finalists use two. Final neural predictions average five refitted seeds. The selected neural family and the selected simpler comparator are chosen separately on development error and fixed before final evaluation. Their comparison can therefore reject the neural structural claim even when the network beats additivity.

Within each development fold, the validation responses choose the best neural stopping epoch and score that checkpoint. Candidate scores average member-specific errors across seeds, rather than scoring a seed-averaged ensemble. Final refits use the median selected epoch across folds and seeds, then average predictions from five seeds before scoring. Development errors are model-selection scores, not unbiased estimates of final ensemble error. For partner cross-validation, predictions from overlapping validation folds are averaged per covered pair before computing each seed's score; uncovered pairs do not contribute.

This bounded search is the initial registered implementation batch, not the roadmap's full capacity expansion. Learning curves use nested development subsets and three subset orders while keeping the full-development-selected hyperparameters. They measure fixed-hyperparameter sensitivity, rather than a subset-retuned sample-efficiency optimum. Additional combination-partner holdouts remove all development pairs incident to specified held targets while preserving their single profiles. Partner folds may overlap in their test pairs and cannot be counted as independent experiments.

3.5 Classical residual kernels

The exploratory genetic kernel family operates directly on observed single-effect profiles. For reference singles $x_u = \delta_u$, the radial kernel is $k_{uv} = \exp[-\|x_u - x_v\|_2^2/(hd_0)]$, where d_0 is the median positive pairwise squared distance between permitted reference singles and $h > 0$ is a selected bandwidth. Two symmetric pair kernels are

$$K_\Sigma(AB, CD) = (k_{AC} + k_{AD} + k_{BC} + k_{BD})/4, \tag{5}$$

$$K_\Pi(AB, CD) = (k_{AC}k_{BD} + k_{AD}k_{BC})/2. \tag{6}$$

The global kernel is the linear Gram matrix of the scaled additive response, divided by the gene count; its per-gene scale is $\max(\sqrt{2}\,\mathrm{sd}(x_{u,g}), 10^{-6})$. Candidates include a shared offset, global correction, the sum kernel, the pair-product kernel, and global-plus-pair mixtures with registered unit or tenfold pair weights. Neither a pair-product form nor a favorable total error is alone evidence of biological interaction learning.

For n training pairs, $H = I - \mathbf{1}\mathbf{1}^\mathsf{T}/n$ centers the training Gram matrix K. With residual matrix R and mean $\bar{r}$, the fitted coefficients are

$$C = (HKH + \alpha I)^{-1}(R - \mathbf{1}\bar{r}^\mathsf{T}), \qquad \widehat{r}(q) = \bar{r} + k_c(q)C. \tag{7}$$

Here $k_c(q)$ uses training kernel means for query centering. The mean residual is an unpenalized intercept; $\alpha > 0$ is selected using development folds. This is closed-form regularized regression, with stored references and coefficients counted in its resource cost.

The contextual amendment uses a separate implementation. It standardizes genes using all explicitly permitted measured single references, including query-context singles, and estimates the reference-distance scale from at most 512 seed-fixed reference profiles. Component kernels are normalized by training diagonal means. Context-conditioned candidates multiply the pair kernel by an RBF of control-derived context features, standardized using source training contexts. A weighted global component and a pooled pair component are optional. These reference allowances make this transductive observed-reference transfer, rather than zero-shot prediction from an unseen context without measured singles. No held-context pair outcome sets reference normalization or fits the source response coefficients.

3.6 Endpoints and uncertainty

For N test pairs and G frozen genes, the primary endpoint and additive skill are

$$\mathrm{MSE} = \frac{1}{NG}\sum_{p,g}(\widehat{\delta}_{p,g} - \delta_{p,g})^2, \qquad \mathrm{skill}_{\mathrm{add}} = 1 - \frac{\mathrm{MSE}_{\mathrm{model}}}{\mathrm{MSE}_{\mathrm{additive}}}. \tag{8}$$

Pairs have equal weight regardless of cell count. Secondary artifacts record absolute error, response correlation, residual alignment, component-informed panel error, and residual geometry. Component panels use single effects and control dispersion where available, without examining held-out doubles. Evaluator-only emergent-response analyses do not define the fitting gene universe.

Paired improvement is baseline MSE minus selected-model MSE. We report 2,000-draw pair bootstrap and target-product-weight bootstrap sensitivities. These are descriptive intervals conditional on estimated profiles and the fitted selection. They do not propagate shared-control uncertainty, the full measurement process, model selection, or independent biological replication. Split-half profiles provide a measurement-reliability diagnostic, not an accuracy ceiling. Ensemble dispersion is not a calibrated confidence interval or a validated abstention policy.

3.7 Synthetic engineering suite

The suite uses 24 targets, 96 genes, and all 276 unordered distinct pairs. Eight generators produce log-scale additivity, count-scale additivity evaluated on a log scale, low-rank interactions, cross-coordinate interactions, higher-rank interactions, a shared nonlinear additive-response distortion, emergent effects in genes without single effects, and noisy singles with a shared control error. Pair responses receive Gaussian observation noise with standard deviation 0.015. Five total random splits and five partner folds are completed for the low-rank generator. These constructed tasks validate implementation behavior and identifiable failure modes; they do not establish biological performance or prevalence of low-rank interactions.

4 Results

4.1 Synthetic recovery depends on the generating structure

Scenario	Selected family	MSE	Skill	Simpler MSE
additive log	bilinear	0.00022	-0.000	0.00022
additive count	bilinear	0.00050	0.471	0.00054
low rank	network	0.00035	0.968	0.00598
cross coordinate	network	0.00026	0.956	0.00240
high rank	bilinear	0.00943	0.657	0.01134
global distortion	network	0.00110	0.941	0.00104
emergent	network	0.00024	0.859	0.00085
noisy controls	network	0.00241	0.842	0.00562

Table 1: Completed synthetic suite. Scores are equal-pair held-out mean squared errors and additive skill as a fraction. These are engineering tasks, not biological validation.

The log-additive negative control has selected MSE 2.22709×10^{-4} and additive skill -0.019%, so residual learning supplies no useful gain. The low-rank and cross-coordinate generators yield skills of 96.76% and 95.63%, respectively,

for development-selected symmetric networks. Their selected errors are 3.48993×10^{-4} and 2.57490×10^{-4}. This demonstrates recovery when the synthetic structure is learnable under the implemented information access.

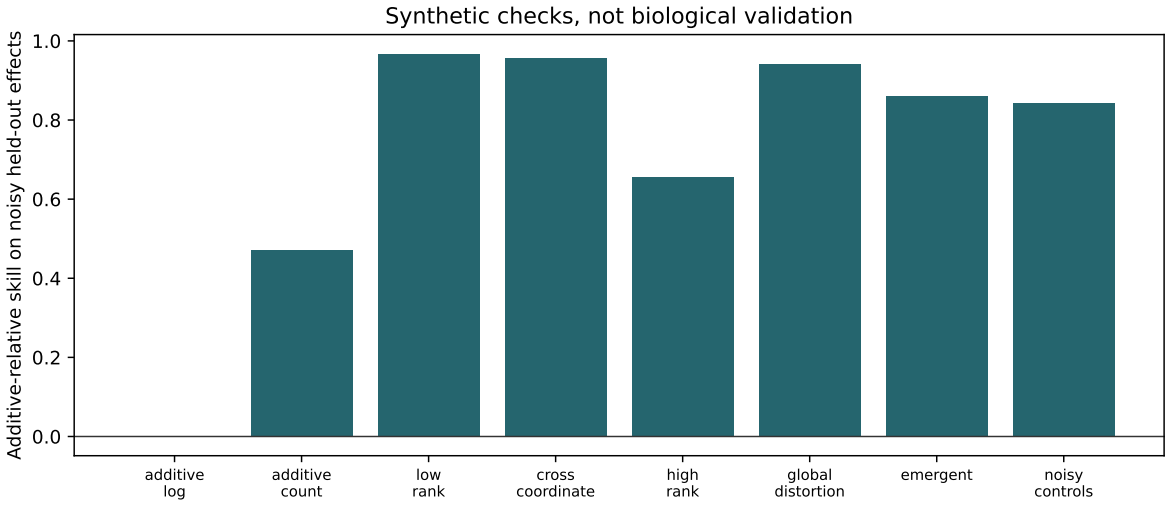


Figure 3: Additive skill across synthetic mechanisms. Positive skill measures predictive improvement, not mechanistic evidence. The shared-distortion case requires comparison to the global correction.

The count-additive generator produces 47.10% skill on the log endpoint, although its underlying count-scale generator has no added interaction. This illustrates scale dependence. In the global-distortion scenario, selected neural skill is 94.13%, but its MSE 0.00110447 exceeds the simpler correction's 0.00104015. A large gain over additivity can therefore coexist with failure of the pair-model claim. The higher-rank scenario selects an unrestricted-output bilinear candidate and yields 65.71% skill; this result does not support a universal low-rank restriction. Emergent and noisy-control scenarios produce 85.95% and 84.24% skill within their constructed universes.

Across the five random splits of the low-rank generator, selected skill ranges from 96.58% to 97.42%. Five partner folds yield 92.97%–96.63% skill, with 86 or 105 test pairs per fold. Generalization survives this harder synthetic split, but variation across held target groups and overlap among test sets limit interpretation. These results cannot substitute for real-data partner holdouts.

4.2 Norman improves over additivity without a clear structural advantage

Dataset/scale	Neural family	Additive MSE	Neural MSE	Simpler MSE	Skill
Norman (log)	network	0.00852	0.00717	0.00727	0.158
Wessels (log)	sum network	0.01394	0.00542	0.00421	0.611
Norman (count)	sum network	1.24605	1.03911	1.25944	0.166
Wessels (count)	sum network	0.26123	0.19684	0.16913	0.246
Chemical (exploratory)	bilinear	0.01381	0.01223	0.00994	0.114

Table 2: Completed primary biological results and available count sensitivity. Architecture and comparator selection use development data only. Skill is a fraction, and errors across expression scales are not directly comparable.

On 38 withheld Norman pairs, the selected rank-two symmetric network has MSE 0.007169187, compared with 0.008515040 for additivity, giving 15.81% skill. Its separately development-selected reduced-rank regression comparator has MSE 0.007268184. The selected network's paired improvement is only 0.0000989974, with descriptive pair-bootstrap interval $[-0.000239746, 0.000393522]$ and target-product-weight interval $[-0.000477775, 0.000691408]$. Both cross zero, so the structural comparison is inconclusive under these conditional summaries.

An ordinary ridge candidate has the lowest observed classical test MSE, 0.006891604, and 19.07% additive skill. This candidate was not the frozen development-selected simpler comparator. Its favorable test score is a descriptive comparison among prespecified evaluated candidates, not grounds to select a replacement after viewing the test. It nevertheless prevents a broad claim that the selected network outperforms the evaluated classical alternatives.

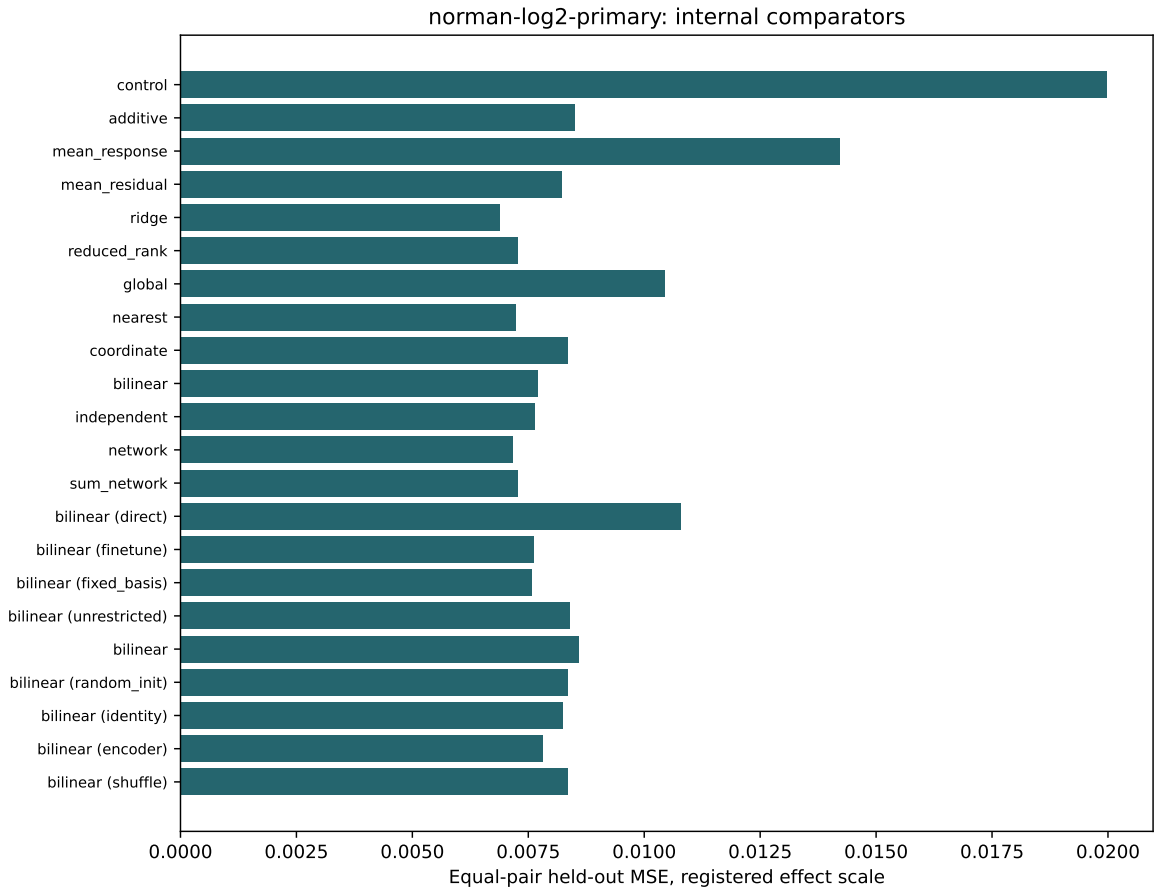


Figure 4: Norman primary internal comparisons. The development-selected neural model and simpler comparator are fixed before test access; the best observed test candidate does not become a new selection.

4.3 Wessels improvement is explained more effectively by a global correction

On 18 withheld Wessels pairs, the selected rank-eight network without multiplicative features has MSE 0.005423558, compared with 0.013936912 for additivity. Its additive skill is 61.08%. The separately development-selected global correction has MSE 0.004214949 and skill 69.76%, so the selected network is worse by 0.001208610 MSE. Pair-specific multiplicative structure has not demonstrated an advantage in this primary experiment.

For the network-minus-global comparison expressed as improvement, the descriptive pair-bootstrap interval is $[-0.00155317, -0.00091829]$, and the target-product-weight interval is $[-0.00189104, -0.00070451]$. Both favor the global correction conditional on these estimated profiles. They are not confirmatory population intervals. The additive comparison is positive under both descriptive schemes, but it cannot distinguish a general correction from pair-specific interaction learning. The selected model's omission of multiplicative features further limits the original architectural hypothesis.

The Wessels count-scale sensitivity also selects the network without multiplicative features. Its MSE is 0.19684419, versus 0.26122700 for additivity and 0.16912624 for the development-selected simpler comparator. Skills are 24.65% and 35.26%, respectively. The negative neural-versus-simpler conclusion persists on this alternative scale, although the magnitudes of count and log errors are not comparable.

Norman count-scale sensitivity selects a rank-eight network without multiplicative features, with MSE 1.03911249 versus additive 1.24604893, giving 16.61% skill. Its development-selected ridge comparator has MSE 1.25943739. The neural improvement over ridge is 0.22032490, but the descriptive pair interval $[-0.06229010, 0.53790781]$ and target-product-weight interval $[-0.26011966, 0.85000429]$ both cross zero. The count-scale target-weight interval for

improvement over additivity also crosses zero. Scale sensitivity therefore supports neither a stable product-feature selection nor a resolved classical-versus-neural advantage.

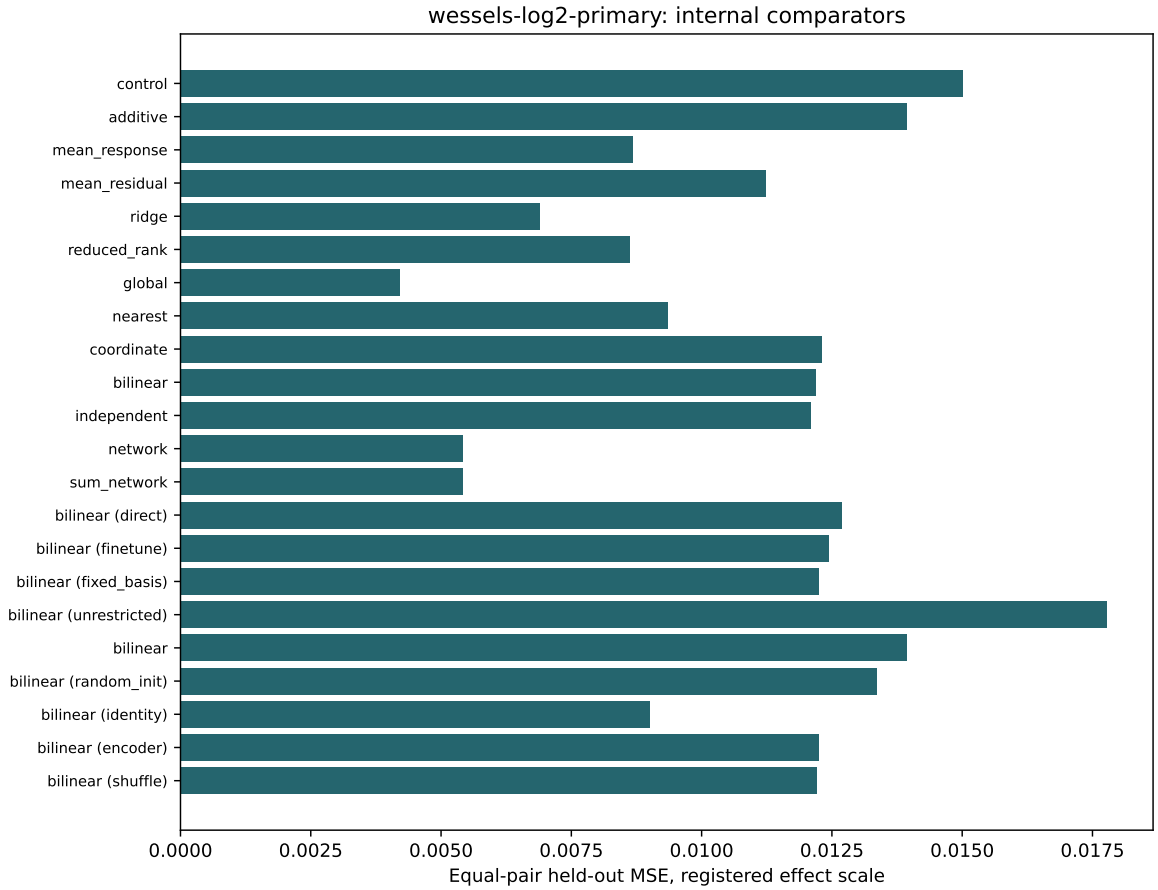


Figure 5: Wessels primary internal comparisons. The global correction outperforms the selected network under the same observed-single input contract.

4.4 Reliability and computational diagnostics

Mean split-half residual cosine is 0.2891 for Norman and 0.0838 for Wessels. These technical partitions indicate substantial instability in the estimated interaction residuals; they are neither independent biological replicates nor an upper bound on prediction accuracy. An evaluator-only, swapped-half emergent panel of 20 genes per pair gives MSE 0.2401 and 0.1268, respectively. This diagnostic covers only the primary frozen gene universe, so it cannot establish recovery of genes excluded during reference-based filtering.

Ensemble-spread versus pair-error Pearson correlations are 0.5262 and 0.5835 on development out-of-fold predictions for Norman and Wessels, and 0.5441 and 0.3124 on their tests. These associations do not calibrate a failure probability, and no risk threshold is selected. Retrieval against the run's held-out responses identifies 35 of 38 Norman pairs and 17 of 18 Wessels pairs, compared with chance rates $1/38$ and $1/18$. Retrieval is sensitive to candidate composition and signal strength, so it supplements expression-error evaluation rather than replacing it. Prespecified best, median, and worst error examples are retained without selecting appealing biological narratives.

Selected five-member ensemble checkpoints occupy 6,079,406 bytes for Norman and 5,439,360 bytes for Wessels. With one CPU thread and 100 warm measurements, median single-pair prediction takes 0.676 and 0.751 milliseconds; 95th percentiles are 0.797 and 1.015 milliseconds. Cold checkpoint load takes 47.0 and 45.7 milliseconds. These retained measurements are saved separately as immutable timing records. They describe this machine and artifact layout, not an end-to-end request service. A separate fresh-process measurement after loading and 100 predictions reports peak resident working sets of 222,375,936 and 223,993,856 bytes, including Python and native-library overhead.

Training peak memory was not measured. Baseline, capacity, fixed-hyperparameter learning-curve, residual-geometry, reliability, and example figures are generated alongside the result artifacts.

4.5 Repeated random and combination-partner holdouts

The biological secondary suite completes four additional random-pair seeds per dataset, giving five random splits including the primary, plus five combination-partner holdouts per dataset. Every split independently repeats development selection and final refitting under the same bounded protocol. Partner holdouts exclude development doubles incident to their held targets but retain those targets' measured singles; inner validation uses analogous partner splits. Gene eligibility and normalization remain fixed from allowed references. Overlapping test pairs and shared profiles make these sensitivity runs correlated, rather than independent experimental replications.

Norman neural additive skill ranges from 0.0820 to 0.1581 across random splits and from 0.0111 to 0.1607 across partner splits. Improvement over the independently development-selected simpler method is positive in one of five random splits and four of five partner splits, with partner differences spanning -0.0000812 to 0.0017721 MSE. This protocol dependence does not establish a general structural advantage. Wessels additive skill ranges from 0.5067 to 0.6161 across random splits and 0.2106 to 0.3220 across partner splits; its simpler comparator outperforms the neural selection in all ten runs. These counts are descriptive, not binomial evidence from independent trials.

Dataset	Protocol	Runs	Neural skill range	Gain over simpler range
Norman	random	5	$[0.082, 0.158]$	$[-0.00065, 0.00010]$
Norman	partner	5	$[0.011, 0.161]$	$[-0.00008, 0.00177]$
Wessels	random	5	$[0.507, 0.616]$	$[-0.00162, -0.00097]$
Wessels	partner	5	$[0.211, 0.322]$	$[-0.00139, -0.00025]$

Table 3: Completed correlated biological split sensitivities. Gain is simpler MSE minus selected neural MSE; ranges are extrema, not confidence intervals.

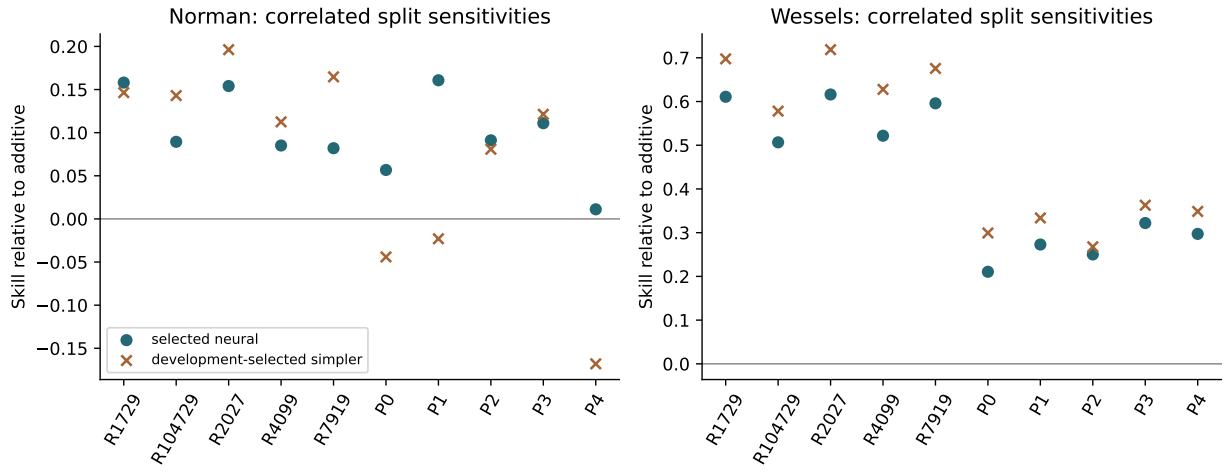


Figure 6: Neural and simpler additive skills across registered biological splits. R denotes random-pair seeds and P partner folds.

4.6 Supplementary adapted public reference

We execute two explicitly adapted SLIM references from pinned upstream revision `5a7e9ade5d0a` (Hu et al., 2026); the full identity and verified source hashes are retained in the method manifest. Its native combination fit is applied to one signed-effect pseudoprofile per condition, retaining the shared gene basis, ridge solve, mean component embedding, and separate single/double biases. We omit the native nonnegative-expression prediction clip because negative effects are valid endpoints. These changes mean the results are adapted references, not reproductions of the published expression protocol. One variant uses 64 supplied STRING dimensions and discloses external information; another substitutes PCA of the allowed observed singles, capped at 16 dimensions.

For each input variant, five development folds select from basis ranks 2, 8, and 16 and ridge penalties 0.1, 1, 10, and 100. Basis fitting uses only singles and that fold's development doubles. These supplementary comparisons were added after the primary results and do not replace the original Opal selection. On Norman, STRING and observed-single variants have MSE 0.00898725 and 0.00782653, respectively, both above the selected Opal MSE. On Wessels their MSE values are 0.00417759 and 0.00418530, both below the selected Opal value and close to the stronger internal global correction.

Dataset	Adapted SLIM input	MSE	Additive skill
Norman	STRING64	0.008987	-0.055
Norman	observed single PCA	0.007827	0.081
Wessels	STRING64	0.004178	0.700
Wessels	observed single PCA	0.004185	0.700

Table 4: Late supplementary signed-effect SLIM adaptations. External STRING information and common-information observed-single PCA are reported separately.

The original Python 3.14 environment lacks native GEARS and CPA dependencies, and CPA requires an older Python runtime. Later isolated Python 3.10 GPU runs execute both methods without changing upstream algorithm source; their aggregate-input results are reported in the context-transfer amendment. PertAdapt's supplied resources pass the retained audit, but its later full-gene float32 forward/backward feasibility probe exhausts RTX 4060 memory, and its distributed launcher requires NCCL unavailable on this Windows build. No PertAdapt biological score is reported. Environment or memory failures document execution scope, not scientific inferiority.

4.7 Extended universe and metric controls

An evaluator-only sensitivity uses all 33,694 original Norman genes and 20,639 Wessels genes. The learned prediction is retained on its frozen schema; outside it, the extension uses observed-single additivity, with no fitted correction. A panel selects the 20 largest absolute residual genes on one technical half and scores the opposite half, then swaps the halves. Outside-schema genes constitute 0.197% of these Norman panel positions and none of the Wessels positions. Selected-extension MSE is 0.240990 and 0.126811, compared with additive values 0.385201 and 0.329619. This sensitivity does not show that a trained extended-universe model would recover emergent responses.

Metric controls include a zero effect, observed additivity, development mean residual, an oracle response as an explicit positive control, and 100 pair permutations of evaluator residuals. These are diagnostics rather than newly eligible candidates. On Wessels, permuted residuals retain mean response Pearson correlation 0.5638 and residual cosine 0.2516 despite MSE 0.020850, above additive MSE 0.013937. The zero-effect prediction also has positive residual cosine 0.6990 while its additive skill is negative. Positive alignment alone can therefore be driven by shared reference geometry and cannot establish useful prediction.

4.8 Conditional technical measurement sensitivity

Two supplementary 200-draw bootstraps resample cells within each condition and matched stratum while retaining QC, gene eligibility, normalization, fitted weights, and method predictions. Each draw uses one matched-control sample per stratum and one single sample per target, reused for every affected pair. This preserves reference sharing in the resampled profiles. It does not propagate parameter fitting, selection, biological replication, or uncertainty from predicting with newly estimated singles.

For improvement over the development-selected simpler prediction, Norman's mean difference is 0.00010518 with a technical profile percentile range $[0.00004388, 0.00016553]$. Wessels' mean is -0.00120615 with range $[-0.00129652, -0.00113008]$. These ranges describe technical response-estimation sensitivity conditional on fixed predictions. Norman's pair-network uncertainty remains unresolved despite the positive technical range; the schemes answer different questions and must not be combined into a confirmatory confidence interval.

4.9 Earlier chemical feasibility does not confirm transfer

Original ComboSciPlex GSE206741 metadata records A549 cells, 24-hour exposure, and 1 micromolar per component, consistent with the primary CPA methods (Lotfollahi et al., 2023). Of 25 pair conditions, only seven have both singles in that experiment, covering six compounds. An explicitly exploratory run uses four development pairs with three-fold cross-validation and three held-out pairs, equal-well profiles, and shared plate controls. Its bounded search contains 22 configurations and omits the genetic ablation batch. Nominal embedding dimension is five.

The exploratory chemical selection is a rank-two bilinear model, with MSE 0.01222957 and 11.44% additive skill. A development-selected mean-residual correction has MSE 0.00994341 and 27.99% skill, so the bilinear selection is worse by 0.00228616. Its additive comparison's descriptive interval crosses zero. Three held-out pairs are insufficient for confirmatory inference, model ranking, or the intended chemical method-transfer claim; these scores document source feasibility and another unfavorable structural comparison. At this historical cutoff the chemical release gate remained unresolved. The later registered EmeraldBay and D-SPIN amendments below provide larger source-qualified tasks, but do not establish positive independent transfer. Unsupported profiles cannot be reconstructed from held-out pairs.

4.10 Coverage and independent transfer

On a common Wessels validation population, expanded development pair coverage improves error by 16.05% in random folds and 4.92% in partner folds; the tested increase in single coverage changes error by about $+0.03\%$ and -0.38%, respectively. These are fixed-population coverage sensitivities, not evidence that pair gains scale indefinitely or that more precise matched references are unnecessary. Full menus and unfavorable architecture, annotation, state, and scale continuations remain in the supplement.

The independent D-SPIN evaluation refits the frozen EmeraldBay-selected procedure on source contexts and withholds two lineages across batches. Frozen-procedure MSE is 0.034025 versus 0.033108 for a source-selected pooled-pair kernel, so the superiority gate fails. A nominal 90% source radius covers 34.3% of held conditions. Query controls and measured singles are allowed, and the held lineages are not independently initiated cultures. This result leaves positive independent transfer and reliable shifted uncertainty unestablished.

5 Discussion

The Wessels result is useful because it separates two claims. Observed-single additivity benefits from correction under this protocol, but the selected neural correction does not surpass the stronger global baseline. A shared expression distortion or broad response calibration can account for predictive gains without evidence that projected pair products capture additional biology. The global synthetic control reproduces precisely this distinction. Its role is diagnostic, rather than proof that the same mechanism generates the biological data.

Output rank and representation quality remain empirical choices. The synthetic unrestricted-output selection shows that a low-rank constraint can be inappropriate, while synthetic low-rank recovery establishes only that the pipeline can exploit a compatible generator. The bounded real-data rank search does not characterize the full spectrum of reproducible biological residuals. Similarly, training-time or parameter-count savings require total resource accounting, including embeddings, single reference profiles, ensemble members, and upstream resources. An inexpensive adapter alone does not establish a small inference system.

The primary endpoint covers the frozen reference-informed universe, but a low overall error can hide failure on emergent genes or poorly measured conditions. Correlation is descriptive and cannot establish accurate residual prediction. Additional endpoints and positive metric controls are required to connect aggregate error to experimental usefulness (Ahlmann-Eltze et al., 2025; Miller et al., 2026). Neither selected correction coefficients nor output directions are validated biological mechanisms.

6 Limitations and completion criteria

This manuscript is an exploratory comparative research preprint, rather than a confirmatory validation study. The later larger chemical pilot, independent-source procedure evaluation and native aggregate public-method adaptations close earlier execution gaps, but positive independent transfer, canonical cell-level comparisons, expanded capacity matching, selection-aware and biological-replicate uncertainty, and biological interpretation remain necessary. Repeated biological splits, count-scale sensitivities, adapted SLIM references, extended-universe diagnostics, technical-cell bootstraps, and a seven-pair chemical feasibility experiment are complete, but do not resolve these gates. The completed reliability diagnostics do not establish a usable abstention policy. The small Wessels held-out set and restricted single coverage particularly limit generalization.

Randomness from optimization is represented by finalist and ensemble seeds, but the reported bootstrap sensitivities do not include the full selection or refitting process. The original technical-cell bootstrap holds predictions fixed, while pair and target-product bootstraps hold estimated profiles fixed. The later technical fitting study refits kernels and coefficients under joint reference draws but fixes model selection, covariance estimation, QC and normalization. Lane and gemgroup matching reduce a source of mismatch without creating biological replication. Cell counts improve

estimates of condition means; they do not increase the number of independent perturbation experiments. Source licensing also requires dataset-specific review before redistribution.

The implemented protocol rejects retrospective replacement of an unfavorable selected model. Further tuning must create a new documented run and preserve the current result. The full research claim will depend on whether improvements survive stronger same-information alternatives and independent contexts. Current evidence supports synthetic recovery, an inconclusive historical Norman structural comparison, and a negative historical Wessels comparison despite gains over additivity. The locally prespecified Stage 1 continuation retains tuned simple corrections; no new pair/state component passes its development gate. A later dose study earns a narrow experimental fixed-curve interpolation route for known compounds and states, while extrapolation fails. Wessels neural models fit inner training labels closely but validate poorly, favoring additional qualified pair/reference evidence over capacity expansion. A metadata-balanced joint development test queries all 54 chemical conditions once but finds no pair-kernel improvement. Independent-culture attribution and sealed globally held pair/context transfer remain missing; their absence does not justify lowering the gate.

Data and code availability

At the current implementation update, version 0.1.2 is a verified private experimental GitHub prerelease from 9 October 2026. Its recorded release and platform receipts apply to their original revisions. The current local development checkout and source-ineligible, scientifically unrun context-contrast successor do not inherit clean-wheel, cross-platform, biological or redistribution certification. The pending-release language in the historical paragraph below describes its manuscript cutoff; no public archive DOI or arXiv submission is claimed.

The repository is prepared at `https://github.com/TheTinyIntelLab/opal`; public release and arXiv submission are pending at this manuscript cutoff. The accompanying source, compact evidence, and Norman and EmeraldBay kernel reference exports are intended for an experimental research release. Original Opal code is MIT; the two dataset-derived reference exports retain CC BY 4.0 terms. Raw measurements, Wessels and D-SPIN weights, and external pretrained resources are excluded from redistribution. Source identifiers, versions, checksums, acquisition adapters, registered protocols, fitted exports, and environment records document the reproduction path. Norman uses Zenodo record 7278143 v1.1; Wessels and the earlier chemical feasibility source use GEO GSE213957 and GSE206741. EmeraldBay uses a pinned Tahoe Bio release (Tahoe Bio, 2026); D-SPIN uses DOI 10.22002/2cjss-wgh69 (Jiang et al., 2026). Evidence tables and publication figures are generated from saved score records without fitting or reopening held-out expression outcomes. The historical baseline passed 117 tests in 12 files; the preserved successor test record passed 139 tests in 18 files; release closeout checks are recorded separately. Software checks and restoration comparisons are engineering evidence; inference restoration does not imply identical CPU/GPU scientific refitting. No public archive DOI or arXiv identifier is claimed at this cutoff.

Ethics declaration

The reported work reuses published cell-line and immune-cell perturbation data and conducts computational experiments. It reports no new human recruitment, animal experiments, or prospective intervention study. Original experimental oversight is described in the source publications. No new ethics approval or exemption is claimed here. The predictions have not been validated for clinical or therapeutic use.

Author contributions, funding, and competing interests

Arjun Vijay Prakash is the sole author and an independent high school student at City Montessori School, with no institutional research affiliation. The work received no funding. The author declares no competing interests. AI assistance is disclosed below; the author is responsible for the final work.

AI assistance disclosure

OpenAI Codex assisted with repository implementation, experimental orchestration, and manuscript drafting. Numerical claims in this draft refer to recorded experiment artifacts rather than generated estimates. AI assistance does not establish authorship, independent replication, or expert validation. The sole author is responsible for the code, analysis, source attribution, and final text. The use of AI assistance is disclosed independently of authorship.

References

Britt Adamson, Thomas M. Norman, Marco Jost, et al. A multiplexed single-cell crispr screening platform enables systematic dissection of the unfolded protein response. *Cell*, 167(7):1867–1882.e21, 2016. doi: 10.1016/j.cell.2016.11.048.

Constantin Ahlmann-Eltze, Wolfgang Huber, and Simon Anders. Deep-learning-based gene perturbation effect prediction does not yet outperform simple linear baselines. *Nature Methods*, 22:1657–1661, 2025. doi: 10.1038/s41592-025-02772-6.

Ding Bai, Le Song, and Eric P. Xing. PertAdapt: unlocking single-cell foundation models for genetic perturbation prediction via condition-sensitive adaptation. *Bioinformatics*, 42(Supplement_1):btag307, 2026. doi: 10.1093/bioinformatics/btag307.

Seowon Chang, Anna Shcherbina, Tal Ashuach, Shahin Mohammadi, Stephanie See, Ninad Ranadive, Emily Fox, and Navpreet Ranu. PertSpectra: Interpretable matrix factorization for predicting functional impact of genetic perturbation experiments. In *Pacific Symposium on Biocomputing*, 2026. URL `https://psb.stanford.edu/psb-online/proceedings/psb26/chang_s.pdf`.

Yiqun Chen and James Zou. GenePT: A simple but effective foundation model for genes and cells built from ChatGPT. *bioRxiv*, 2023. doi: 10.1101/2023.10.16.562533. Preprint.

Dewei Hu, Marc Pielies Avellí, Lars Juhl Jensen, and Simon Rasmussen. SLIM: A small linear model with STRING embeddings for single-cell genetic perturbation prediction. *bioRxiv*, 2026. doi: 10.64898/2026.08.07.743481. URL `https://www.biorxiv.org/content/10.64898/2026.08.07.743481v1`. Preprint; citation verified from author repository.

Jialong Jiang, Sisi Chen, et al. D-SPIN constructs regulatory network models from scRNA-seq that reveal organizing principles of perturbation response. *Cell*, 189(13):4159–4184.e35, 2026. doi: 10.1016/j.cell.2026.04.028.

Arthur Liberzon, C. Birger, H. Thorvaldsdóttir, M. Ghandi, Jill P. Mesirov, and Pablo Tamayo. The molecular signatures database hallmark gene set collection. *Cell Systems*, 2015. doi: 10.1016/j.cels.2015.12.004.

Zeming Lin et al. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379:1123–1130, 2023. doi: 10.1126/science.ade2574.

Russell Littman. Data distribution for PRESAGE (Perturbation Response EStimation with Aggregated Gene Embeddings), 2025. URL `https://zenodo.org/records/15587986`. Genentech copyright; semantic derivatives subject to author noncommercial terms.

Mohammad Lotfollahi, Anna Klimovskaia Susmelj, Carlo De Donno, Leon Hetzel, Yuge Ji, Ignacio L. Ibarra, Sanjay R. Srivatsan, Mohsen Naghipourfar, Riza M. Daza, Beth Martin, Jay Shendure, Jose L. McFaline-Figueroa, Pierre Boyeau, F. Alexander Wolf, Nafissa Yakubova, Stephan Günnemann, Cole Trapnell, David Lopez-Paz, and Fabian J. Theis. Predicting cellular responses to complex perturbations in high-throughput screens. *Molecular Systems Biology*, 19:e11517, 2023. doi: 10.15252/msb.202211517.

Henry E. Miller, Gabriel M. Mejia, Francis J. A. Leblanc, Brendan Swain, Bo Wang, and Lucas Paulo de Lima Camillo. Deep learning perturbation models can outperform baselines on calibrated metrics. *Nature Biotechnology*, 2026. doi: 10.1038/s41587-026-03307-w.

Thomas M. Norman, Max A. Horlbeck, Joseph M. Replogle, Alex Y. Ge, Albert Xu, Marco Jost, Luke A. Gilbert, and Jonathan S. Weissman. Exploring genetic interaction manifolds constructed from rich single-cell phenotypes. *Science*, 365:786–793, 2019. doi: 10.1126/science.aax4438.

Stefan Peidli, Tessa D. Green, Ciyue Shen, et al. scperturb: harmonized single-cell perturbation data. *Nature Methods*, 21:531–540, 2024. doi: 10.1038/s41592-023-02144-y.

Arjun Vijay Prakash. Opal registered dose development and prospective precision sensitivity, 2026a. Local research successor 0.1.2.dev2. Protocol and missingness amendment in configs/experiments/dose-design-20261007.json and dose-design-amendment-20261007.json; evidence in results/dose-design-20261007. Exposed development and hypothetical design sensitivity.

Arjun Vijay Prakash. Opal publisher annotation identity and shared-target development, 2026b. Local successor 0.1.2.dev3. Protocol configs/experiments/identity-reference-20261007.json; exposed development and metadata-only provenance audits in results/identity-reference-20261007.

Arjun Vijay Prakash. Opal global chemical-pair and whole-state development, 2026c. Local successor 0.1.2.dev4. Protocol configs/experiments/joint-development-20261007.json; exposed development evidence in results/joint-development-20261007. Independent confirmation remains unavailable.

Arjun Vijay Prakash. Opal matched scale and source-selected panel development sensitivity, 2026d. Successor 0.1.2.dev5. Protocol configs/experiments/panel-scale-20261007.json; exposed development results and diagnostic subsets in results/panel-scale-20261007. No independent confirmation.

Arjun Vijay Prakash. Opal registered stage 1 successor development protocol and evidence, 2026e. Local research successor 0.1.2.dev1. Protocols in configs/experiments/stage1-20261007.json and its explicit information supplement; results in results/stage1-20261007. Exposed development, not independent confirmation.

Jason Qin, Hans-Hermann Wessels, Carlos Fernandez-Granda, and Yuhan Hao. Active learning for efficient discovery of optimal combinatorial perturbations. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 50363–50383, 2025. URL `https://proceedings.mlr.press/v267/qin25g.html`.

Jingang Qu, David Holzmüller, Gaël Varoquaux, and Marine Le Morvan. TabICLv2: A better, faster, scalable, and open tabular foundation model. arXiv:2602.11139, 2026. URL https://arxiv.org/abs/2602.11139.

Joseph M. Replogle et al. Mapping information-rich genotype–phenotype landscapes with genome-scale Perturb-seq. *Cell*, 2022. doi: 10.1016/j.cell.2022.05.013. URL https://pmc.ncbi.nlm.nih.gov/articles/PMC9380471/.

Yusuf Roohani, Kexin Huang, and Jure Leskovec. Predicting transcriptional outcomes of novel multigene perturbations with GEARS. *Nature Biotechnology*, 42:927–935, 2024. doi: 10.1038/s41587-023-01905-6.

Glenn Shafer and Vladimir Vovk. A tutorial on conformal prediction. *Journal of Machine Learning Research*, 9(12):371–421, 2008. URL https://jmlr.org/papers/v9/shafer08a.html.

Sanjay R. Srivatsan, José L. McFaline-Figueroa, Vijay Ramani, et al. Massively multiplex chemical transcriptomics at single-cell resolution. *Science*, 367:45–51, 2020. doi: 10.1126/science.aax6234.

Tahoe Bio. EmeraldBay, 2026. URL https://huggingface.co/datasets/tahoebio/EmeraldBay. Dataset revision f2a0be6b02f731553657f0115c345b20bb020ede; CC BY 4.0; accessed 6 October 2026.

Frederik Wenkel, Wilson Tu, Cassandra Masschelein, et al. TxPert: using multiple knowledge graphs for prediction of transcriptomic perturbation effects. *Nature Biotechnology*, 2026. doi: 10.1038/s41587-026-03113-4.

Hans-Hermann Wessels, Alejandro Méndez-Mancilla, Yuhan Hao, Efthymia Papalexi, William M. Mauck, Lu Lu, John A. Morris, Eleni P. Mimitou, Peter Smibert, Neville E. Sanjana, and Rahul Satija. Efficient combinatorial targeting of RNA transcripts in single cells with Cas13 RNA Perturb-seq. *Nature Methods*, 20:86–94, 2023. doi: 10.1038/s41592-022-01705-x.

A Exploratory residual-kernel development

After inspecting the completed primary and secondary findings, we investigated a closed-form symmetric residual kernel procedure. The candidate set contains a shared offset, reference-scaled linear global correction, RBF summed-single and symmetric pair-product kernels, and mixtures of global and pair-product corrections. Inner development folds select among 67 configurations. This procedure permits a shared correction to win; the original primary selection restricted the selected method to a neural family. Historical model selections and test results remain unchanged.

On each genetic development set, three seeds define five outer random folds and five outer partner folds, with three inner folds of the corresponding protocol per outer fit. An expanded linear comparator selects among 25 configurations comprising a shared offset and global, ridge, and rank-four reduced-rank models at eight penalties. The kernel procedure reduces outer-validation MSE relative to this comparator by 1.78–2.70% for Norman random holdouts, 1.13–3.38% for Norman partner holdouts, 2.50–3.98% for Wessels random holdouts, and 0.83–1.00% for Wessels partner holdouts. These are development-only point estimates across dependent resamplings, not independent confirmation. The search spaces and budgets differ, so these results do not establish a capacity-matched structural advantage.

Partner folds overlap; the primary nested score weights validation occurrences equally. Averaging repeated predictions changes the estimand and can reverse Wessels' comparison against the reference-scaled shared global control. No universal advantage over linear methods is inferred. The historical test findings motivated this amendment, and no new historical evaluator responses were loaded. Safe candidate checkpoints reproduce the saved unscored query predictions. Native frontier comparisons, independent confirmation, and novelty remain unresolved. Evidence is recorded separately under `results/improvements/`; it does not replace the frozen primary experiments.

A separately recorded exploratory check adds 45 frozen scFoundation positional-gene-table kernel candidates to the measured-single search, with an identically searched shuffled-gene-table control. These vectors are checkpoint gene positions, not contextual backbone outputs or a native PertAdapt replication. Norman retains the same measured-single candidates in every outer fold. On Wessels, the real table's small random-holdout improvement is matched or exceeded by shuffling gene identities, and its partner score is slightly worse. This does not support an identity-specific benefit; the measured-single candidate remains the exported model. Pretraining-data overlap is unknown.

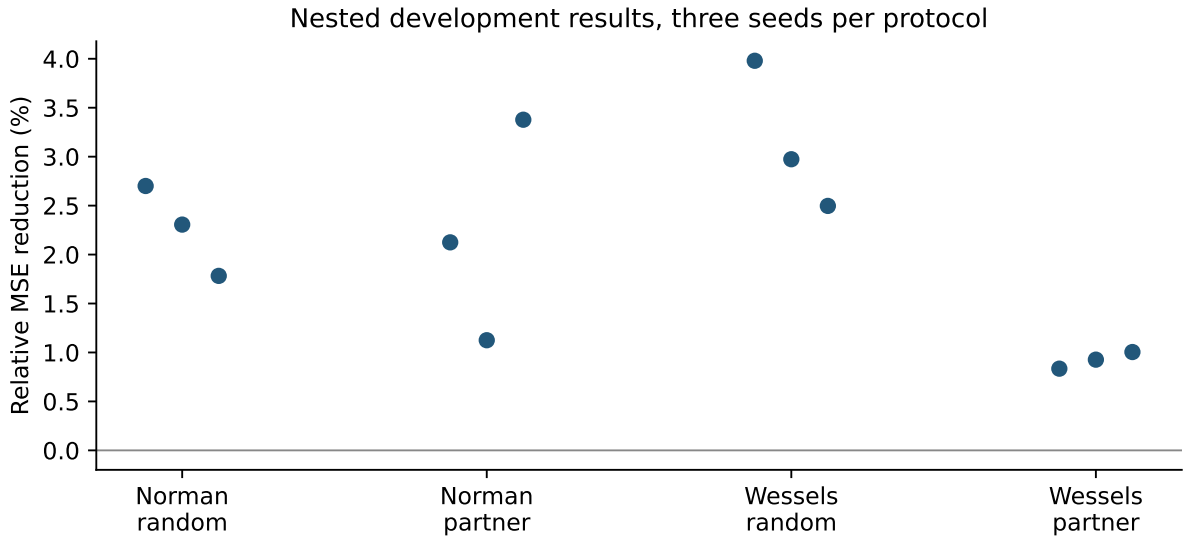


Figure 7: Exploratory nested-development improvement relative to a separately tuned expanded linear search. Each point is one split seed; the points are dependent and are not confidence intervals.

B Architecture and coverage development

All studies in this section were designed after the historical outcomes were known and use development responses only. They are exploratory. Nested five-outer/three-inner random and partner folds select configurations, while historical evaluator responses remain inaccessible. The primary nested score weights validation occurrences equally; overlapping partner folds and repeated seeds are dependent.

B.1 Reference uncertainty, priors and structured outputs

Six implementation priorities are tested: technical reference uncertainty, multiple biological priors, structured output directions, a bounded signal-sensitive loss, independent single-screen transfer, and strong classical experts. Sixteen joint technical bootstrap draws reuse the same matched-control draw within each stratum and the same single-target draw across affected pairs. QC, gene eligibility and normalization remain fixed. Reference reliability shrinks measured-single features and supplies jointly sampled training augmentation. It does not estimate biological variance.

The source bank combines measured-single PCA, a fixed random projection of the supplied GO graph, Hallmark pathway membership (Liberzon et al., 2015), and, in the external-prior tracks, frozen GenePT and ESM2 embeddings from the verified PRESAGE author archive (Littman, 2025; Chen and Zou, 2023; Lin et al., 2023). Missing sources are masked. Neural heads fuse sources and form symmetric summed, absolute-difference and product features. Output directions comprise training-fold residual PCA, pathway-local single-reference directions, regularized gene-local GO directions, and optional external directions. The gene-space loss is computed exactly through the output-basis quadratic. Its bounded signal weights incorporate training response magnitude and technical variance; this is an Opal objective, not an implementation of another study's published weighted metric.

Twelve neural ablations, each with two mixture choices, are compared with the 67-configuration residual kernel and a separately tuned prior-feature ridge. Ablations include uncertainty, fusion, pooling, structured modules, shared output rank and signal weighting. The kernel is selected in all forty competitive outer folds across the original and expanded datasets. This selection permits a simple expert to win. It does not establish a matched-search-budget architecture advantage. The separately nested measurement GP and training-only residual retrieval are worse than the kernel on the original datasets. GP conditional latent intervals omit query measurement variance, full reference covariance and selection uncertainty, so their empirical coverage is diagnostic only.

Dataset	Fold	Kernel	Neural selection	GP	TabICLv2
Norman	random	0.006292	0.006809	0.007012	0.007443
Norman	partner	0.007181	0.007538	0.007878	0.008529
Wessels original	random	0.005032	0.006990	0.008983	0.007458
Wessels original	partner	0.009868	0.011249	0.012159	0.012613
Wessels expanded	random	0.006471	0.010523	0.011666	0.012228
Wessels expanded	partner	0.014467	0.016805	0.017884	0.018610
Wessels protected	random	0.007360	0.011987	–	–
Wessels protected	partner	0.016063	0.018735	–	–

Table 5: Development-only architecture comparisons. Every entry uses the same dataset-specific validation population and weights validation occurrences equally. The neural entry is inner-selected. TabICLv2 uses a fixed CPU configuration, not a native perturbation benchmark. Missing entries indicate that a comparator was not run on that track.

A frozen TabICLv2 regressor (Qu et al., 2026) predicts eight training-fold residual-PC scores from the same symmetric reference/prior features, using four estimators and a pinned, checksum-verified public checkpoint. Its error exceeds the kernel on all six original/broad-expanded dataset-protocol comparisons. Expanded partner folds emitted Yeo–Johnson overflow warnings, although retained predictions were finite. This fixed configuration does not characterize the best tuned foundation-model performance, and pretraining contamination cannot be excluded.

B.2 Measurement-aware kernels and identity controls

Heteroscedastic kernels fit a generalized-least-squares shared correction with an unpenalized intercept. Each training pair's technical residual variance determines a bounded diagonal noise factor; no validation variance enters the fitting likelihood. Inner folds choose noise strengths zero, one-half or one over the 67 kernel candidates. The unweighted search has 67 candidates, while real-variance and permuted-variance searches each have 201. The matched placebo permutes variance values within training folds only. Shared bootstrap draws preserve dependence when estimating variances; this first track uses their diagonal approximation. The subsequent shared-covariance track below tests off-diagonal dependence explicitly.

GenePT, ESM2, GO and Hallmark kernels are also tested against an identically searched shuffled-gene-identity bank. Norman random folds and both expanded random tracks retain the measured-only baseline. Original Wessels random-fold MSE improves from 0.00503184 to 0.00500368 with real priors, but shuffling also improves it to 0.00501316. Norman partner error worsens slightly with real priors. Small improvements on the expanded partner tracks do not establish a general identity-specific benefit.

Dataset	Fold	Real noise gain (%)	Permuted gain (%)
Norman	random	1.27 to 2.54	-0.06 to 0.19
Norman	partner	-0.58 to 1.13	0.00 to 0.00
Wessels original	random	0.21 to 1.05	-0.08 to 0.10
Wessels original	partner	-0.07 to 0.19	-0.01 to 0.00
Wessels protected	random	0.92 to 1.24	-0.05 to 0.02
Wessels protected	partner	-0.40 to 0.05	-0.01 to -0.00

Table 6: MSE reduction versus the unweighted kernel over three dependent split seeds. Ranges are extrema, not confidence intervals. Real-variance and permuted controls have the same search budget.

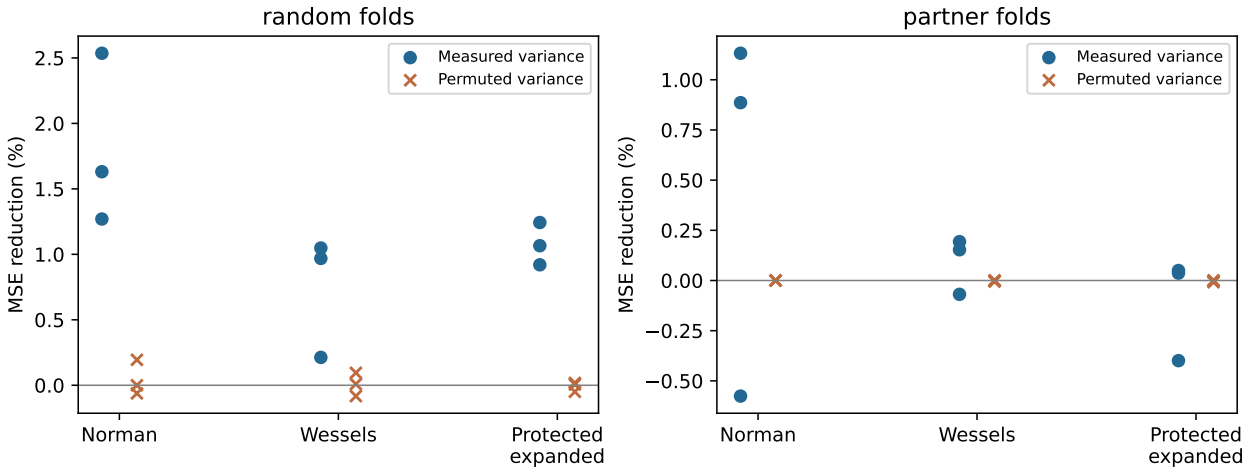


Figure 8: Each point represents a development split seed. Measurement weighting consistently helps random folds in these runs; partner gains are mixed and smaller.

B.3 Independent representation transfer

The verified PRESAGE archive contains semantic embeddings, splits and ranked DEG lists rather than the required expression matrices. We therefore separately download the scPerturb release of the Replogle K562 essential-gene CRISPRi screen (Replogle et al., 2022) and verify its publisher MD5. It contains 310,385 cells, 8,563 output genes and 2,057 single-target conditions. Of these, 2,002 targets map to the supplied GO dictionary. A seed-fixed 1,601/401 training/validation target split trains a 16-to-32-to-16 encoder for 200 updates against training-target response PCs. Its normalized validation-PC MSE is 0.87670, versus 0.88021 for an alpha-ten ridge and 0.90792 for the training mean. These are source-task development scores.

External effects are mean released-X differences against controls. This scale is not equated with Opal's count-derived log2 pseudobulk scale. Only encoder initialization and aligned response directions transfer to the separately fitted current-context pair decoder; external combination labels are not used. Encoder-only, decoder-only and combined transfer are compared with a matched no-transfer head and linear ordinary/weighted objectives. The first transfer run was interrupted after unscaled semantic feature products destabilized its linear controls. That evidence is retained. A second run retained poor fixed-update linear optimization and was also interrupted. Both records remain available. The final linear controls standardize features using training pairs only, preserve scaler buffers in safe checkpoints, and solve the convex gene-space quadratic by converged conjugate gradient. Their frozen-reference fitting omits stochastic reference augmentation; the ordinary and weighted linear objectives otherwise share features and regularization. Results from the corrected track are generated below.

Encoder initialization modestly improves the matched Norman random-fold neural head, but its partner score is worse. None of the encoder, decoder or combined-transfer heads beats the kernel on any dataset-protocol aggregate. The weighted linear objective has higher primary MSE than its ordinary counterpart on all eight aggregates. These results limit the transfer and signal-objective hypotheses under the tested representation and regularization.

The unbounded convex controls have particularly poor error on some expanded partner folds. A separate late exploratory control clips standardized features to magnitude eight in both fitting and inference, retaining the same con-

Dataset	Fold	No transfer	Encoder	Decoder	Both	Linear MSE	Linear weighted
Norman	random	0.008892	0.008549	0.008895	0.008613	0.008281	0.008517
Norman	partner	0.010484	0.010785	0.011279	0.010869	0.009534	0.009743
Wessels expanded	random	0.010658	0.011229	0.010621	0.011237	0.023141	0.028550
Wessels expanded	partner	0.019227	0.022255	0.019632	0.023492	0.415128	0.479667
Wessels protected	random	0.012031	0.012413	0.012093	0.012512	0.027270	0.031739
Wessels protected	partner	0.021088	0.022368	0.021576	0.023081	2.346815	2.903400
Wessels original	random	0.007193	0.007259	0.007775	0.007398	0.006618	0.006678
Wessels original	partner	0.016089	0.015757	0.016792	0.016447	0.010943	0.011011

Table 7: Corrected transfer and convex objective controls on development folds. Linear heads solve a regularized gene-space quadratic on frozen standardized features; neural heads use joint reference augmentation. These are different fitting procedures, and none uses historical query responses.

verged ordinary and weighted quadratics. It is designed after the unfavorable observations, and its results do not replace them. Exact feature-extreme magnitudes are omitted because the compact release records do not supply a traceable numerical source for them.

Dataset	Fold	Kernel	Linear MSE	Bounded MSE	Bounded weighted
Norman	random	0.006292	0.008281	0.008190	0.008413
Norman	partner	0.007181	0.009534	0.008810	0.008947
Wessels expanded	random	0.006471	0.023141	0.010730	0.011126
Wessels expanded	partner	0.014467	0.415128	0.017244	0.017678
Wessels protected	random	0.007360	0.027270	0.012413	0.012865
Wessels protected	partner	0.016063	2.346815	0.019325	0.019807
Wessels original	random	0.005032	0.006618	0.006615	0.006675
Wessels original	partner	0.009868	0.010943	0.010921	0.010985

Table 8: Late exploratory bounded-feature controls after observing expanded-fold extrapolation. The standardized feature map clips at magnitude eight during training and prediction; the regularized gene-space objectives are solved to convergence. This is a separate amendment, not a replacement of the unfavorable unbounded controls.

B.4 Expanded coverage under a common validation population

Lowering Wessels reference eligibility to fifteen cells yields 28 singles and 158 matched pairs, compared with the original 16 singles and 58 pairs, while retaining the 10,590-gene schema. A broad exploratory repartition assigns 110 pairs to development and 48 to query, moving twelve historical query identities into development. A second expansion preserves all original development/query identity roles and assigns only new pairs. Its development set contains no original query identity. Nevertheless, thirty of its query identities had already appeared in the broad development track; it is not independent confirmation.

A controlled comparison scores the same forty original development pairs under the same outer folds and gene panel. Additional protected development pairs may enter training, but extra pairs containing excluded inner or outer partner targets are removed. Extra singles alone change random error by only 0.03% and worsen partner error by 0.38%. Extra eligible pair labels reduce random error by 16.05% and partner error by 4.92%. Adding measurement weighting gives 18.73% and 4.91% reductions. This supports a coverage benefit for this exploratory population, not a new architecture mechanism or independent biological result. Across three dependent split seeds, extra pair labels reduce random-fold error by 16.05–17.60% and partner error by 4.00–4.99%. With measurement weighting the ranges are 18.73–19.54% and 4.08–4.99%, respectively. These ranges are split sensitivities, not confidence intervals. A stricter sensitivity removes extra pairs touching any validation component and gives smaller partner gains; both policies are retained.

The common `main.py frontier` entry point prepares resources, repeats registered studies, validates exported inference, generates evidence tables and builds this multi-file manuscript. Large licensed inputs, prediction banks and derived weights remain local and excluded from source-control commits. Safe model exports include reference banks, numerical states, immutable gene order, checksums and contextual input guards. These implementation and empirical gains do not satisfy an independent frontier-quality claim.

Fold	Original	Extra singles	Extra pairs	Pairs + noise
random	0.005032	0.005030	0.004224	0.004089
partner	0.009868	0.009905	0.009382	0.009383

Table 9: Common-population Wessels nested-development MSE. Original query identities are excluded from all training. Expanded response histories are known, so these scores are exploratory.

B.5 Shared covariance, fitting sensitivity and retuned data efficiency

A subsequent exploratory track retains the full technical residual covariance across development pairs. Sixteen joint cell-bootstrap draws reuse shared single and control samples, and their residual covariance is averaged over the frozen output genes. The likelihood uses only each fitting fold's covariance block. Generalized least squares retains an unpenalized shared intercept; inner folds choose variance strength and off-diagonal shrinkage. The unweighted, diagonal and correlated searches contain 67, 201 and 469 candidates respectively. A jointly permuted training covariance has the same 469-candidate budget as the correlated family. No query responses enter these studies.

Dataset	Split	Diagonal	Correlated	Permuted
Norman	random	1.26 to 1.48	0.97 to 1.47	-0.20 to 0.00
Norman	partner	0.86 to 1.29	0.77 to 1.30	0.00 to 0.15
Wessels original	random	0.36 to 1.30	0.36 to 1.06	-0.03 to 0.00
Wessels original	partner	-0.42 to -0.07	-0.42 to -0.07	-0.03 to 0.02
Wessels protected	random	-0.06 to 0.90	0.29 to 0.90	-0.04 to 0.02
Wessels protected	partner	-0.39 to -0.16	-0.96 to -0.28	-0.00 to 0.00

Table 10: Percentage MSE reduction relative to the unweighted kernel over three additional development split seeds. Ranges are extrema of dependent resamplings, not confidence intervals. Negative reductions indicate worse error.

The correlated family improves Norman random error by 0.97–1.47%, while diagonal weighting improves it by 1.26–1.48%. Original Wessels random gains are 0.36–1.06% for correlated noise and 0.36–1.30% for diagonal noise. Both families worsen original and protected Wessels partner aggregates in these three seeds. Correlated protected random gains are 0.29–0.90%. Thus the extra covariance parameters do not establish a consistent gain over the diagonal likelihood, and the additional split histories expose sensitivity in earlier positive findings.

For every correlated outer fit, each of the sixteen technical reference draws recomputes the reference-derived kernel and refits response coefficients. Predictions are scored against the point-estimated outer validation response. Recorded 5th, 50th and 95th error percentiles and gene-level prediction standard deviations quantify fitting/reference sensitivity conditional on the selected configuration, covariance, QC and normalizer. They do not include selection uncertainty or biological replication and are not confidence intervals for population performance.

Training subsets of one-quarter, one-half, three-quarters and full outer training size independently retune each kernel family using fitting-only inner folds. Three nested subset orders share the same outer validation set; full size is fitted once. Infeasible inner partitions are retained without response-dependent redraws. A fixed-full-size-hyperparameter kernel provides a comparator. A first curve run reordered the full-size pair list, changing its inner partition; it is retained. The corrected run uses canonical full-size order, and verifies identical fixed and retuned unweighted full-size predictions. Larger subsets improve error in these runs, while retuning itself has mixed effects. No independent sample-efficiency advantage follows from the correlated covariance extension.

A separate diagnostic divides each outer fitting population into fitting and calibration pairs. Only fitting responses select hyperparameters and fit the diagonal kernel. A condition-level gene-RMSE radius uses the finite-sample rank $\lceil (n+1)c \rceil$ for requested coverage c; ranks above the calibration sample count have no finite radius. Conformal coverage statements require appropriate exchangeability assumptions (Shafer and Vovk, 2008); these dependent biological pairs do not establish them. We report empirical outer coverage without implementing a calibrated abstention policy.

Original Wessels random folds have seven calibration pairs, so a nominal 90% finite radius is unavailable on every fold. Protected Wessels partner folds attain only 83.7% empirical coverage at nominal 90%, and original Wessels partner folds attain 84.7% over their four finite-radius folds. These results do not support a usable guaranteed-risk release. The common `main.py ceiling` route reproduces this extension, and `main.py frontier` includes it in the full architecture build.

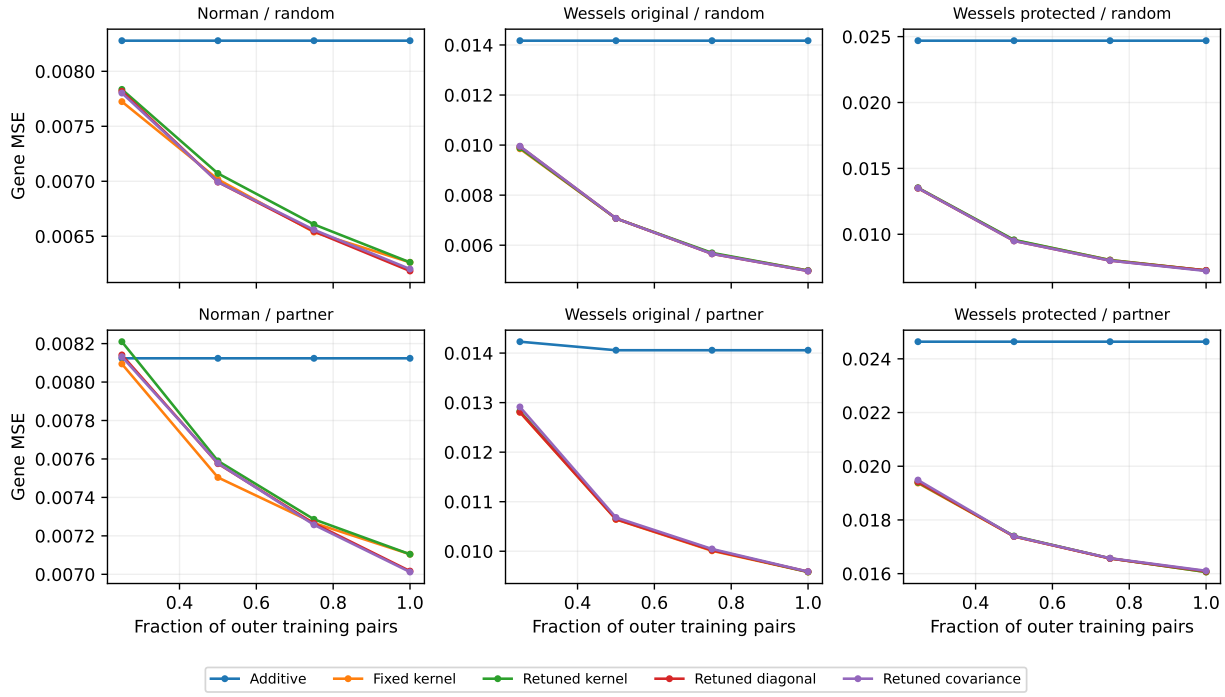


Figure 9: Subset-retuned development learning curves with a fixed-hyperparameter kernel comparator. Subset orders and partner validation occurrences are dependent. Infeasible small-subset fits can change the scored occurrence population.

Dataset	Split	Nominal	Finite folds	Observed	Mean radius
Norman	random	80	5/5	81.8	0.0982
Norman	random	90	5/5	94.3	0.1270
Norman	partner	80	5/5	87.0	0.1195
Norman	partner	90	5/5	94.4	0.1333
Wessels original	random	80	5/5	85.0	0.1005
Wessels original	random	90	0/5	–	–
Wessels original	partner	80	5/5	71.8	0.1105
Wessels original	partner	90	4/5	84.7	0.1239
Wessels protected	random	80	5/5	76.4	0.1023
Wessels protected	random	90	5/5	90.0	0.1374
Wessels protected	partner	80	5/5	76.7	0.1604
Wessels protected	partner	90	5/5	83.7	0.1708

Table 11: Isolated error-radius diagnostic on one development split seed. Nominal and observed coverages are percentages; mean radii are gene-RMSE units. Observed coverage counts only folds with finite radii and is descriptive.

C Context-transfer and native-method amendment

C.1 Contracts, sources and selection

This amendment preserves the historical tests and uses new registered datasets and output directories. Reference banks identify biological context, study, modality, assay and exact intervention/reagent/dose/endpoint. Availability flags distinguish absent profiles from measured zero effects. Fitting combinations, grouped calibration responses, query identities and evaluator responses are physically separate. Measured query controls and singles are allowed; combination predictions are written before calibration and held-response evaluation. Six residual families distinguish a shared offset, global correction, pooled pair or sum kernels, and their context-weighted variants. Source-only context folds select configurations; the kernels are symmetric positive semidefinite and include an unpenalized mean residual.

EmeraldBay is pinned to dataset revision `f2a0be6b02f731553657f0115c345b20bb020ede` (Tahoe Bio, 2026). Six checksum-verified shards contain 94,746 cells. Condition coverage, requiring at least twenty cells per matched control, single and pair, selects eight eligible lines before response-based comparisons. Seventeen exact-dose pair conditions yield 99 context/pair rows: 54 fitting, 18 calibration and 27 evaluation rows, withholding NCI-H1573 and SW1417. The 1,024-gene panel uses vehicle variation only. Unlike the historical transform-after-condition-mean track, these chemical adapters normalize each cell's counts to 10,000, apply log2 transformation and then average. Raw MSE magnitudes are not comparable across these tasks. The EmeraldBay pilot is exposed and exploratory; repeating it creates no new independent evidence.

The independent D-SPIN dosage-combination source contains 173,412 cells (Jiang et al., 2026). Publisher MD5 and recorded SHA256 verify its count matrix. Stimulated controls and exact-dose matching retain 55 pair identities and 273 pair/context rows across five published lineages and two batches. The frozen vocabulary intersection has 1,011 genes. Myeloid and NK are withheld across both batches (108 evaluation rows); inner selection withholds whole source lineages, preventing batch copies from crossing a selection fold. Calibration has eleven pair groups. Published lineage labels are post-treatment, and repeated dose-grid rows or batches do not establish biological replication.

C.2 Observed-single transfer and uncertainty

The EmeraldBay-selected pooled-sum kernel has alpha 0.01 and bandwidth four. It improves over additivity, but equals its development-selected simple reference, giving no additional skill over that stronger alternative. Four pilot calibration groups are insufficient for a finite nominal-90% rank. Before D-SPIN held-response access, this procedure is frozen and refitted on D-SPIN source contexts; EmeraldBay fitted weights are not pooled across assays. The independent source test is complete, but its positive confirmation gate fails: frozen-procedure MSE is 2.8% worse than the source-selected pooled-pair reference, rather than the registered minimum 5% improvement. Later source-selected family scores do not replace the unfavorable frozen result.

The D-SPIN group-max calibration radius is 0.165342 and covers only 34.3% of held conditions. This is an empirical source radius without an out-of-context coverage guarantee. Reference-distance selective-risk diagnostics do not establish a usable abstention policy. These negative outcomes leave both positive procedure transfer and reliable uncertainty unresolved.

C.3 Native GPU methods and attribution

Native CPA uses upstream revision `fbd7c0250edc23eff003a10c99655579c53afd63` and native Gaussian training on allowed aggregate control, single and fitting-combination observations. Source-context selection chooses width 128, latent dimension 32 and 120 epochs; final checkpoint selection uses only allowed reference rows. Its chemical pilot MSE is 8.8% below the selected kernel. This retrospective aggregate-input adaptation is not the canonical CPA cell-level benchmark. Export restoration without combination outcomes reproduces all 27 predictions exactly.

Native GEARS uses revision `f374e43e197b295016d80395d7a54ddb81cc6769`, its architecture and fourth-power plus direction loss, with hidden width sixteen and forty epochs on the registered genetic development folds. Positive-cosine nearest-three graphs derive only from allowed single profiles instead of external ontology graphs; observations are aggregate means. GEARS beats the pooled-sum kernel on Wessels partner folds (0.008419 versus 0.010053 MSE), loses on Wessels random folds (0.007094 versus 0.005031), and loses to the kernels on both Norman protocols. These adapted, fixed-native-versus-tuned-kernel comparisons establish no universal winner or state-of-the-art claim. Five partner folds overlap; their per-occurrence mean differs from scoring averaged repeated predictions.

Retuned kernel attribution compares all six families, fixed branches, shuffled singles and matched reference-removal controls without reopening historical evaluator outcomes. The family differences do not establish an interaction mechanism. Acquisition comparisons share fixed validation populations, tune using initial labels only and compare uniform,

Population / contract	Method	MSE
EmeraldBay	Observed additive	0.21023854
EmeraldBay	Selected pooled-sum kernel	0.17221753
EmeraldBay	Native CPA, aggregate inputs	0.15709786
D-SPIN	Observed additive	0.03334998
D-SPIN	Source-selected pooled-pair kernel	0.03310752
D-SPIN	Frozen kernel procedure	0.03402466
EmeraldBay, masked singles	Estimated additive	0.15722541
EmeraldBay, masked singles	Estimated singles plus residual	0.09282916
Norman development	Native GEARS, random	0.01119871
Norman development	Native GEARS, partner	0.01170730
Wessels development	Native GEARS, random	0.00709371
Wessels development	Native GEARS, partner	0.00841881

Table 12: Saved-score next-cycle evidence. Rows are comparable only within the same population and input contract. Chemical rows use log-transformed per-cell aggregate means; genetic development rows use historical prepared profiles. Native methods are disclosed aggregate adaptations. Masked-single rows use a separate reference bank. Partner-fold means count each validation occurrence; overlapping folds are not independent replicates.

target-balanced, diversity and kernel-variance policies over three dependent seeds. Diversity helps one Norman budget but worsens a Wessels budget; no universal prospective policy is established. Twelve saved ceiling metadata audits pass, documenting consistency rather than new biological validation.

C.4 Estimated singles and reproducible exports

Missing-single inference has a distinct experimental contract. All query singles are masked from both the imputer and residual reference bank. Only measured instances of an exact intervention/dose in source contexts can supply an estimate; absent source identities trigger abstention and availability flags are never relabeled. Source-only single validation selects a control-feature kernel, improving over the target mean by only 0.5%, below the registered 5% complexity gate. This pilot does not establish a material contribution that would justify adding a learned representation.

All 27 pilot queries are supported. Total MSE is 0.092829 versus estimated-additive 0.157225, but main-effect MSE 0.165359 and residual MSE 0.183338 combine with cross term -0.255868. The resulting cancellation prevents interpreting the favorable total as reliable main-effect transfer; it does not rank this masked-reference protocol against the observed-single track. No calibrated interval is provided.

Checkpoint state version three persists fitted gene scaling and reference-distance normalization, avoiding restoration differences from numerical-thread roundoff. Observed and estimated predictions restored from an installed wheel outside the checkout agree within 7×10^{-15} without importing PyTorch. At the historical context-cycle cutoff, the software suite passed 104 tests, covering contracts, kernel properties, split isolation and provenance. Compact evidence and generated manuscript tables carry source hashes; data, prediction arrays and model archives remain local. Completed cycle-v7 reads saved records, while interrupted disk-full cycle-v6 is preserved. Engineering verification does not turn unfavorable scientific gates into a validated release.

D Baseline-only state-response development

This build tests main-effect transfer before expanding the architecture. Matched-control-relative effects, source-only inner selection and whole-biological-group outer folds define the task. Query treatment outcomes never fit control PCA, target means or response bases. The source control panels permit explicitly declared query-control gene selection; fitted state encodings remain source-only.

The tested learned response model combines a source-average intervention anchor with a projected source-response head. Source-only control PCA represents the baseline; exact intervention one-hot identity, source single-effect features, and numeric dose are separate input paths. Numeric dose is the only descriptor used here, rather than a chemical structure representation. Both the state width and response rank are 16. Two registered training settings use 80 epochs with regularization 0.01 and 160 epochs with regularization 0.1, each with seeds 41 and 73. Simple state references and the learned heads share source intervention means and permitted query-control gene-panel adaptation, so their comparison tests added representation value within this information contract.

The corrected scPerturb 1.4 sci-Plex 3 release is publisher-MD5 verified (Srivatsan et al., 2020; Peidli et al., 2024). Its full matrix has 799,317 rows, including unassigned cells and a separate 72-hour A549 block. Restricting to 24-hour human-cell-line conditions, 24 shared exact 1,000-nM intervention identities yield 144 single condition means across 24 line/plate strata and three biological line groups. All plates and replicates of a line remain on one side of each holdout. The 512 human-gene panel uses controls only; count-library normalization uses human genes. Wells receive equal weight, and each effect subtracts its own line/plate-matched vehicle. Plate count is not independent state replication.

Source	Target mean	Selected simple	Learned state
EmeraldBay	0.064758	0.064680	0.073678
sci-Plex 3	0.074913	0.074906	0.091868

Table 13: Nested main-effect context development. Each source has its own scale/vocabulary and fixed evaluation population; MSE is not compared between sources. Neither simple-state nor learned-representation comparison passes the registered 5% gate.

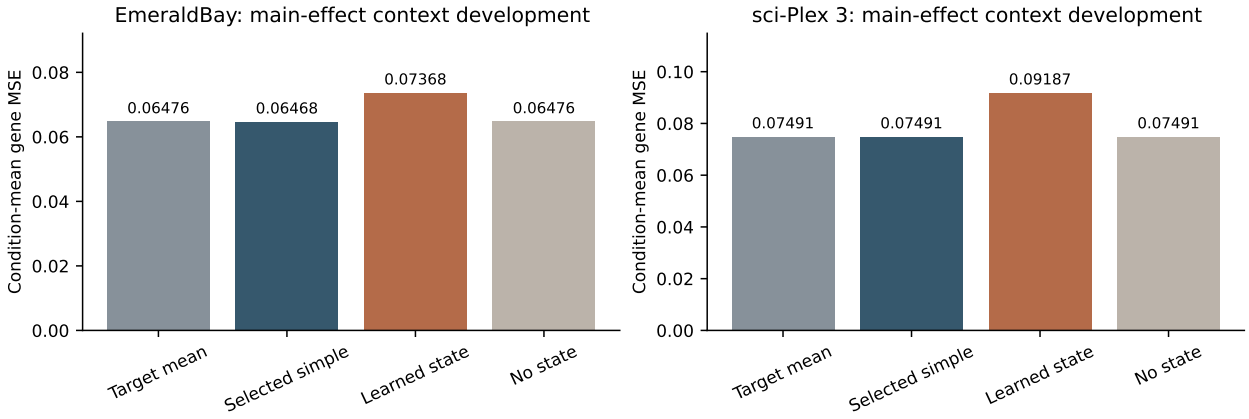


Figure 10: Source-average anchors remain strong. Removed-state and shuffled-input controls use the same supervision and output anchor; they are encoder ablations, not unseen-intervention methods.

State-conditioned simple references improve on target averaging by 0.121% on EmeraldBay and 0.010% on sci-Plex 3. The learned candidates are respectively 13.91% and 22.65% worse than selected simple references. These negative development results close the scaling/routing gate for this build; three sci-Plex lines do not support a broad state-generalization conclusion.

All selected sci-Plex 3 interventions use 1,000 nM, so shuffling the numeric dose descriptor is a no-op and yields identical saved learned scores. This control is uninformative about descriptor specificity. These results reject adoption or scaling of the tested state build; they do not establish that all learned state representations, biological context information, or context-dependent interactions are unhelpful.

Nominal sci-Plex duplicate-well profiles have mean replicate-difference MSE 0.063736, compared with a mean selected-simple improvement of 0.00000718. The replicate contrast includes plate and biological variation, shared controls and estimation error; it is a measurement-sensitivity comparison, without an identified biological accuracy ceiling or independent-culture precision claim.

D.1 State-dependent residual and training-stage comparisons

Only EmeraldBay source-development pair outcomes enter this new experiment; historical pair calibration and evaluator outcomes are not reopened. Inner kernel selection masks query singles and uses the same estimated-single route as outer evaluation. Symmetric sum/difference and product heads train measured-single residuals with the main model fixed. Joint variants retain source-single supervision and independent held-single diagnostics. Frozen and fine-tuned encoders use two residual seeds and equal pair epochs; removing neural single pretraining retains the source-average anchor and bases.

Model	Total	Main	Residual	Cross term
Estimated additive	0.18978	0.18380	0.19237	-0.18639
Selected kernel	0.09461	0.18380	0.17046	-0.25966
Staged state	0.11239	0.18380	0.16674	-0.23815
Frozen-encoder joint	0.10163	0.17534	0.15817	-0.23188
Fine-tuned joint	0.10148	0.17519	0.15755	-0.23126
No neural pretraining	0.09516	0.16712	0.15721	-0.22917

Table 14: Pair/context development on one fixed population. Total error equals main plus residual error plus the cross term. Large negative cross terms prevent interpreting good total error as accurate interaction recovery.

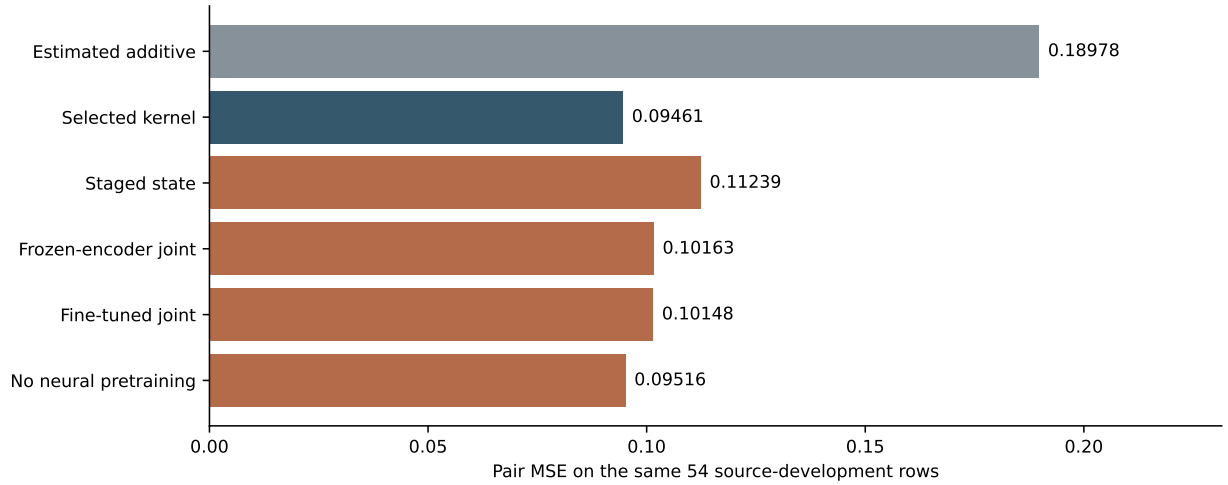


Figure 11: State heads improve their pooled neural counterparts, but none beats the information-matched selected residual kernel. No winner is selected retrospectively from these outer scores.

D.2 Reliability, resources and remaining evidence

A risk diagnostic fits separately from two-group calibration and two-group evaluation. At nominal 90%, two calibration groups cannot produce a finite group-rank radius. Accept-none and accept-all comparisons remain visible; the registered accepted-risk/coverage criterion also fails. This is shared-study development, not fresh risk confirmation. D-SPIN remains a failed transfer of the earlier source radius. Inference returns unavailable intervals and diagnostic support rather than a confidence percentage.

Three subset orders retune simple references at each context budget with fixed query conditions. These dependent retrospective curves and hypothetical condition-assay costs do not establish prospective acquisition value. A new uncertainty-based acquisition extension is withheld after failed gates. Native CPA/GEARS results remain information-scoped historical comparisons; full-gene PertAdapt remains unscored. A separately identified 256-token synthetic adapter probe passes float32, float16 and bfloat16 forward/backward at batches one and two, without proving full-gene memory fit. A Linux/NCCL system with full gene alignment and an initial 48-GB float32 profiling budget is a documented profiling target rather than an established full-gene resource requirement.

Two-seed portable NumPy exports preserve normalization, gene order, exact-condition records, supporting reference identities and all heads. At the state-experiment cutoff, the suite passed 115 tests. Python and CLI routes validate compatible previously unregistered baselines, pair swapping and one-time baseline reconstruction outside the checkout in an isolated installation. Engineering completion does not establish independent confirmation or reliable inference. A replicated, randomized shared-intervention screen with separate sealed confirmation and risk-calibration blocks remains required.

E Locally prespecified Stage 1 development

The local prespecification and evidence archive are identified separately from historical findings (Prakash, 2026e).

The successor preserves commit `84de7e3` as the historical preparation. All reused outcomes are exposed development; no independent confirmation is claimed. Feature preprocessing and model selection use only source rows and declared measured references. A two-seed ensemble is the neural selection unit, with inner-only stopping, 16 configurations per shared/pair family, 120 updates, and unchanged original 80-update schedule controls.

Source / split	Candidate	MSE	Gain (%)
emeraldbay-pilot-v1/context	Selected simple	0.150322	0.00
emeraldbay-pilot-v1/context	context	0.151126	-0.53
emeraldbay-pilot-v1/context	identity	0.144215	4.06
emeraldbay-pilot-v1/context	pair_product	0.152077	-1.17
emeraldbay-pilot-v1/context	mixed	0.150199	0.08
norman-log2/partner	Selected simple	0.007493	0.00
norman-log2/partner	identity	0.007530	-0.50
norman-log2/partner	pair_product	0.007815	-4.30
norman-log2/partner	mixed	0.007750	-3.43
norman-log2/random	Selected simple	0.006564	0.00
norman-log2/random	identity	0.006444	1.82
norman-log2/random	pair_product	0.006860	-4.52
norman-log2/random	mixed	0.006859	-4.50
norman-log2/random	neural_shared	0.007404	-12.79
norman-log2/random	neural_pair	0.007407	-12.85
wessels-log2/partner	Selected simple	0.010506	0.00
wessels-log2/partner	identity	0.010538	-0.31
wessels-log2/partner	pair_product	0.013510	-28.60
wessels-log2/partner	mixed	0.013511	-28.61
wessels-log2/random	Selected simple	0.005470	0.00
wessels-log2/random	identity	0.005797	-5.98
wessels-log2/random	pair_product	0.005531	-1.12
wessels-log2/random	mixed	0.005476	-0.11
wessels-log2/random	neural_shared	0.006332	-15.77
wessels-log2/random	neural_pair	0.006761	-23.61
emeraldbay-pilot-v1/main	Selected simple	0.064643	0.00
emeraldbay-pilot-v1/main	learned	0.071363	-10.40
sciplex-state-v1/main	Selected simple	0.074945	0.00
sciplex-state-v1/main	learned	0.086414	-15.30

Table 15: New exposed development. Gains are against inner-selected simple competitors on the same information contract. Baseline-only main prediction is separate from observed-single pair correction. Original frozen runs supply this table and its figures; the changed-environment recheck remains a named sensitivity. Original frozen runs supply this table and its figures; the changed-environment recheck remains a named sensitivity. Full paired uncertainty and every control are in the saved comparison CSV.

The initial materiality rule requires five percent relative improvement, with a positive paired diagnostic lower bound for continuation. These fixed-profile state/target resampling intervals do not establish independent biological precision. One-hot identity renaming is mathematically invariant; an explicitly registered supplementary row permutation breaks alignment. Numeric dose descriptors do not test chemical identity.

Disjoint technical-cell halves fit shared correction with whole shared-reference matrices retained. Opposite-half residual diagnostics and main/residual/cross-term accounting are measurement sensitivities, not independent cultures or biological interaction recovery. Baseline-only pair ratios against observed-single references are excluded from rankings. Original results remain intact.

Independent joint transfer remains unevaluable. It requires globally held pair identities, multiple sealed pretreatment states, independently initiated cultures with matched controls and singles, and declared query-reference adaptation. Reliability, acquisition and scaling remain conditional.

E.1 Registered dose transfer and prospective precision design

A metadata-qualified extension uses 24 lexicographically selected compounds at four exact doses across A549, K562 and MCF7, at 24 hours. All outcomes are exposed development. An initial plate-specific run stopped on missingness before scoring; an explicit amendment pools matched-control-subtracted plate means within each line. Four global dose holdouts retain 72 profiles each and 512 control-selected genes. Nested source-dose selection compares zero,

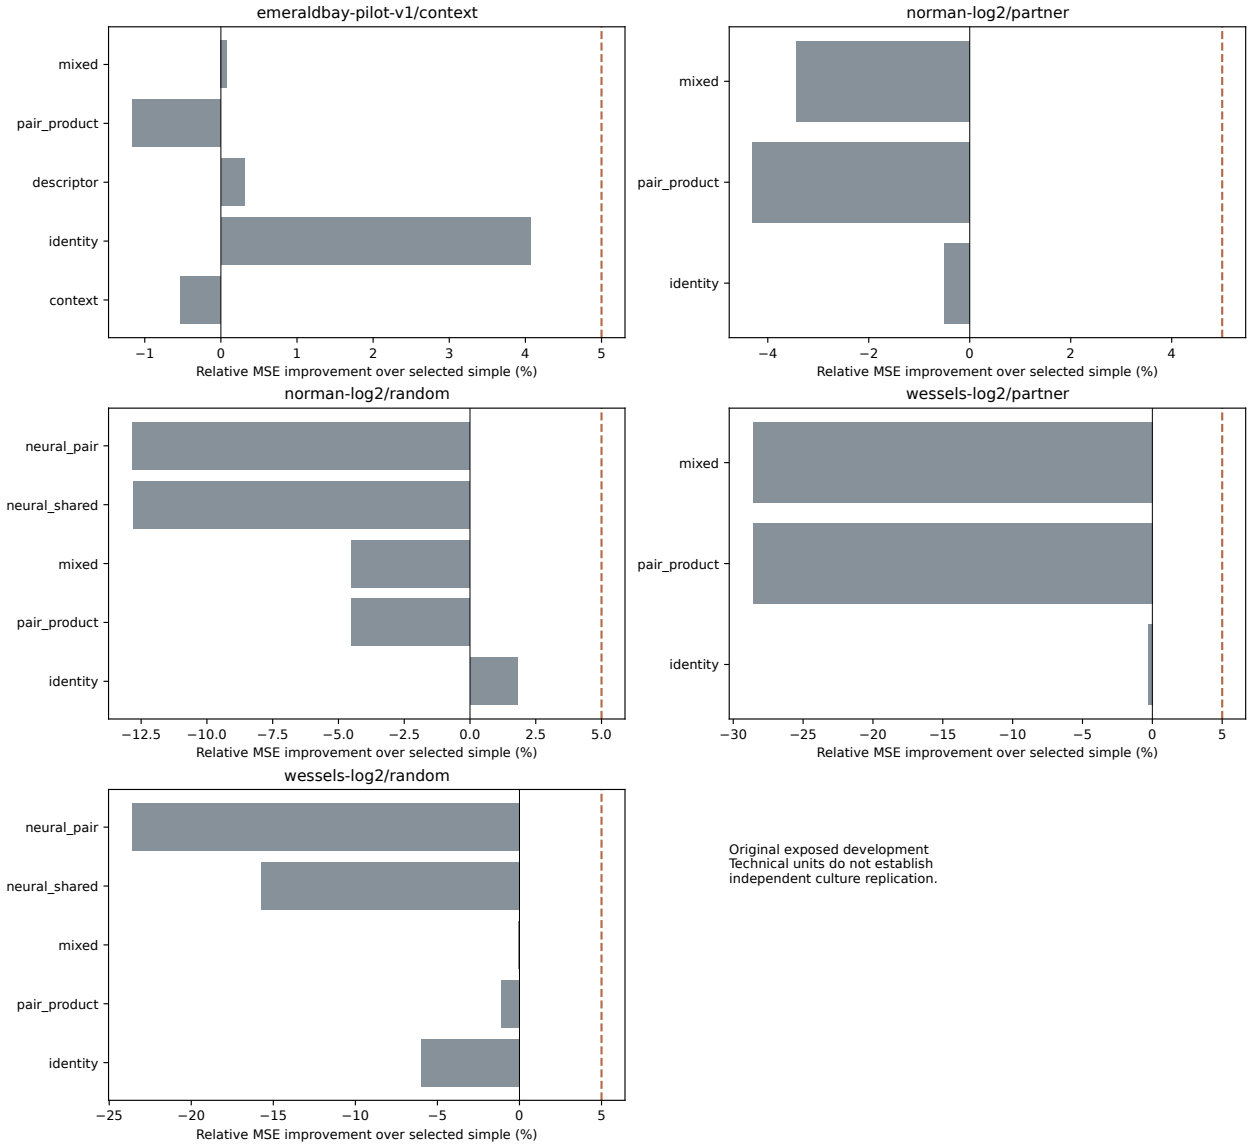


Figure 12: Information, pair attribution and neural changes against retained simple competitors. The dashed line is the locally prespecified development materiality screen.

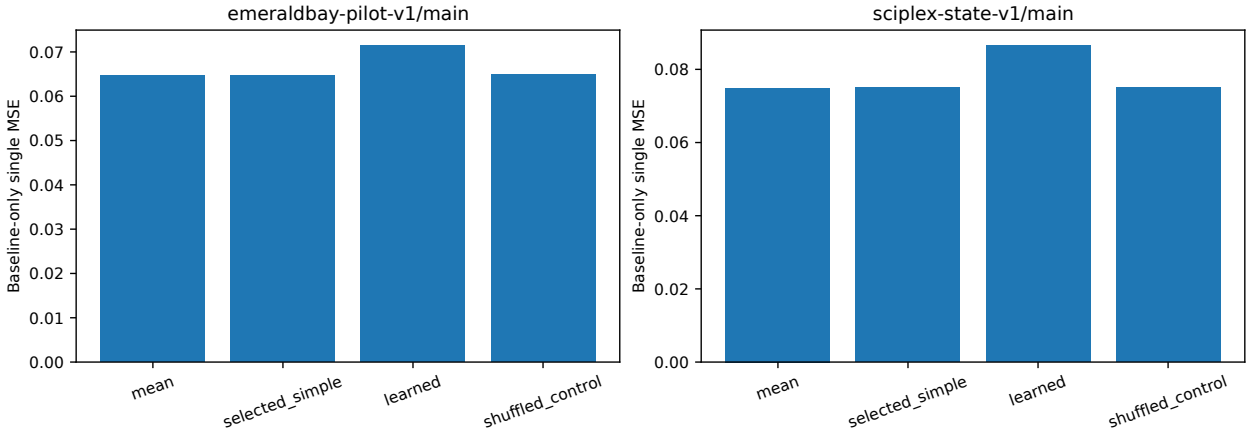


Figure 13: Separate baseline-only main prediction with fixed source identity support. Shuffled-control results are fixed-selection sensitivities.

mean and nearest-dose references with 30 fixed ridge/kernel curves. Query controls may select the gene panel; held-dose effects never enter fitting, scaling, selection or pretraining (Prakash, 2026a).

Dose (nM)	Reference MSE	Smooth MSE	Gain	Paired reduction interval
10	0.037674	0.050863	-35.01%	[-0.022821, -0.000754]
100	0.028909	0.026544	8.18%	[0.000883, 0.004364]
1000	0.052405	0.028459	45.69%	[0.005018, 0.047922]
10000	0.060991	0.068715	-12.66%	[-0.021095, 0.009771]

Both interpolation doses pass the 5% development screen with positive conditional compound-bootstrap lower bounds; both extrapolation doses fail. Fitting takes 1.89 CPU seconds with final working set 200.0 MB. Bootstrap intervals condition on shared controls and three cell lines, excluding biological and selection uncertainty. An experimental known-compound/known-state interpolation reference preserves the selected source-dose curves and rejects extrapolation, unknown identities and endpoint mismatches; restored scored predictions agree within 0. No pair/context or chemical identity transfer claim follows. A separate prospective sensitivity calculation evaluates 320 hypothetical

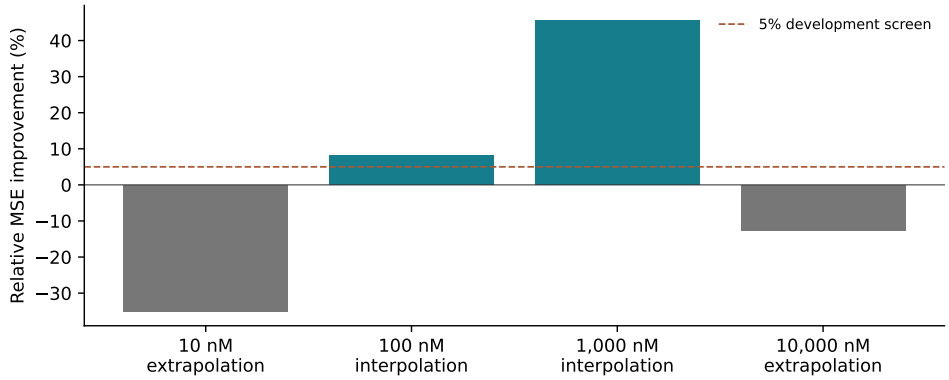


Figure 14: Source-selected smooth dose curves improve interior doses and fail extrapolation on exposed sci-Plex development.

designs with 2,000 Gaussian draws each. Crossed state, culture, target, pair and shared-control effects remain correlated; cell noise alone decreases with cell count. Assumed variance is not estimated from technical halves. The known-variance lower bound is optimistic and does not establish a powered sample size. Four development states with three independently initiated cultures per state form a variance pilot; broad transfer requires multiple sealed states and globally held pairs. The pilot must estimate biological covariance and mean precision before an empirical power calculation.

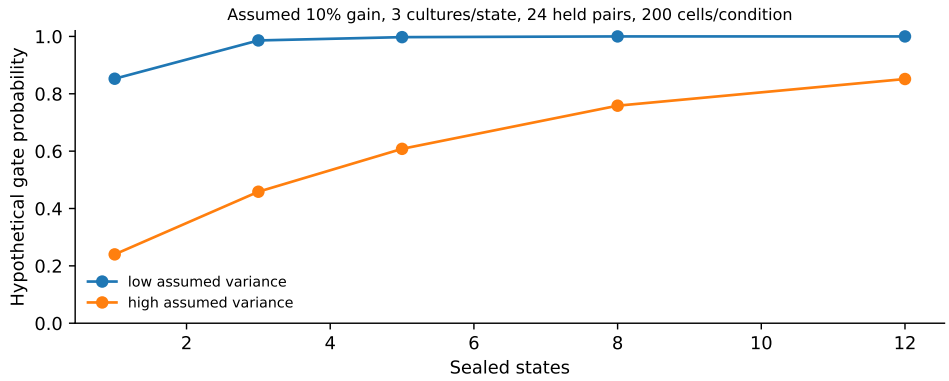


Figure 15: Hypothetical design sensitivity under declared low and high variance assumptions. These are simulated design probabilities, not empirical confidence in Opal.

E.2 Publisher annotation transfer and recovered provenance

A subsequent metadata qualification tests inherited publisher target/pathway annotations from the corrected CC BY 4.0 release. Four identity folds globally exclude every held compound response at all doses and states. A separate four-fold shared-target track takes transitive connected components of comma-separated target classes before balancing folds, preventing a multi-target compound from bridging source and held target labels. Source-only categorical vocabularies retain explicit unknown/missing masks. The common population has 288 profiles from 24 compounds, three known lines and four doses. These annotations are permitted external information, not chemical structures or independently verified mechanisms (Prakash, 2026b).

Track	Reference MSE	Annotation MSE	Gain	Paired reduction interval
Identity	0.047102	0.046110	2.11%	[-0.001893, 0.003979]
Shared target	0.047503	0.049028	-3.21%	[-0.006883, 0.003309]

Neither track passes the 5% screen with a positive conditional lower bound. Shuffled annotations have MSE 0.047574 and 0.049367, respectively. Each fold selects among nine simple/class-mean references, 24 ridge/kernel annotation maps and the same 24-map shuffled search, using whole source identity/component folds only. Both tracks together use 2.44 CPU seconds, with completion working set at most 97.5 MB. Whole-compound bootstrap draws condition on shared controls and known lines; biological, refitting and multiplicity uncertainty are excluded. No annotation encoder or unseen-compound production route earns advancement. A separate primary-metadata audit aligns all 799,317 cell

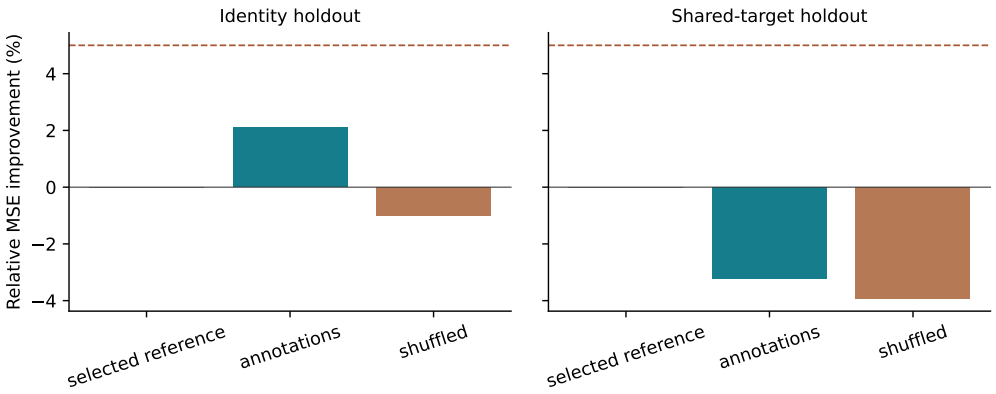


Figure 16: Publisher annotations fail both identity and shared-target development gates; the dashed line marks the registered 5% screen.

identities with no condition or annotation mismatches, without reading expression outcomes. It recovers one original

sample identifier, two replicate labels per line, and 48 line/culture control strata. Original culture labels alias existing assigned plate strata, supplying no independently initiated culture evidence beyond the randomized treatment wells already documented. The count matrix is not refitted. Pair responses, multiple sealed states, biological covariance for empirical power, and independent risk evaluation remain absent. Recovery of metadata nomenclature does not create fresh confirmation.

E.3 Globally held chemical pairs and whole-state development

A separately registered metadata-balanced partition closes an exposed-development coverage gap without reopening the historical evaluator. Seven canonical chemical pairs, defined by unordered names across all doses, are greedily balanced into two global panels and crossed with six whole cell lines. Twelve outer folds query each of the 54 fitting conditions once; guard rows are unused. Source sizes range from 19 to 26. Inner selection also holds a whole state and global chemical pair, retains only preregistered metadata-feasible folds, and uses no outcome-driven redraw. Exact-dose measured query singles and the legacy query-control gene-panel adaptation are explicitly permitted (Prakash, 2026c).

Correction	Total MSE	Gain	Paired reduction interval
selected simple	0.175076	0.00%	[0.000000, 0.000000]
additive	0.192365	-9.88%	[-0.040297, 0.008862]
mean residual	0.201447	-15.06%	[-0.047645, -0.011049]
global correction	0.174662	0.24%	[0.000000, 0.002411]
single sum	0.175103	-0.02%	[-0.004653, 0.004813]
pair product	0.176599	-0.87%	[-0.010852, 0.005217]
mixed	0.175353	-0.16%	[-0.006766, 0.006065]

The source-selected simple reference chooses mean residual, tuned global correction or summed-single kernels by inner MSE. Pair-product and mixed kernels worsen total error by 0.87% and 0.16%; neither passes the registered 5% materiality gate. Both intervals span zero. The registered 2,000 crossed state/intervention-weight bootstrap draws have 1,965 nonempty draws; their intervals condition on shared reference profiles and selection, exclude biological/refitting uncertainty, and remain diagnostic. The fixed search takes 2.16 elapsed seconds on one CPU thread, with 276.4 MB completion RSS, not a measured peak.

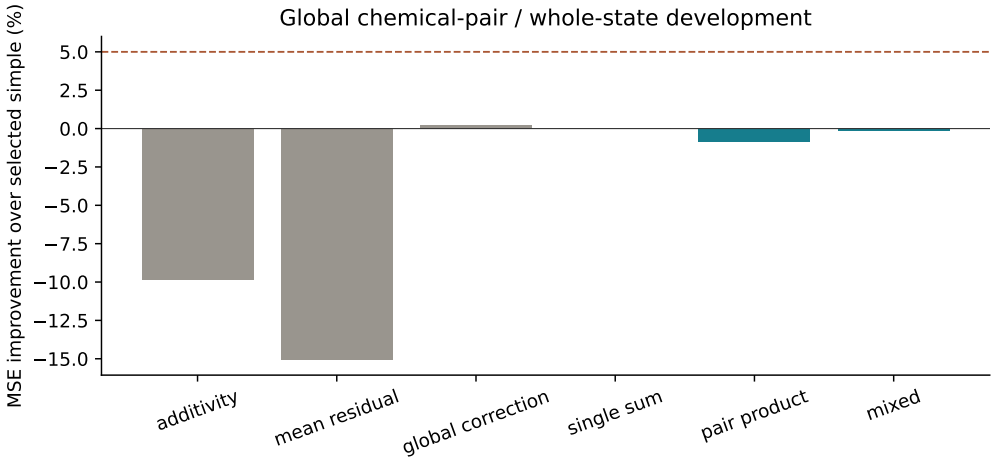


Figure 17: Pair-specific kernels fail the joint development gate against a source-selected strong simple reference. The dashed line marks the registered 5% screen.

Metadata-balanced joint development is feasible here, superseding the earlier single-row lexicographic feasibility check as a development assessment. All outcomes remain exposed. This result does not supply independent joint confirmation, new independently initiated cultures, or enough sealed states to estimate biological uncertainty. No pair head, learned state encoder, risk policy, acquisition or scaling component earns advancement.

E.4 Matched scale sensitivity and source-selected reporting panels

A further registered development supplement addresses scale and reporting gaps without enlarging the neural menu. Metadata qualification verifies that count/log2 banks share raw-source hashes, genes, target/pair identities, original development/query roles, normalizers and QC. Count models reuse the exact three random and three partner outer folds per source, with the original nested source partitions. The fixed menu retains additivity, mean residual, 24 tuned global ridge/RBF settings and 18 settings each for summed-single, pair-product and mixed kernels. The simple reference is source-selected, and neither historical evaluator nor calibration outcomes are accessed (Prakash, 2026d).

Count-scale track	Simple MSE	Pair MSE	Mixed MSE	Pair / mixed gain
norman random	1.099782	1.054091	0.998685	4.15 / 9.19%
norman partner	1.846124	1.676695	1.704124	9.18 / 7.69%
wessels random	0.187987	0.198708	0.195647	-5.70 / -4.07%
wessels partner	0.242347	0.269093	0.267866	-11.04 / -10.53%

Norman count-scale gains are positive, unlike its log2 results, so direction is scale-sensitive. The primary gates still fail: Norman random mixed correction gains 9.19%, with paired MSE-reduction interval $[-0.011700, 0.273082]$; Norman partner pair correction gains 9.18%, with interval $[-0.113894, 0.496702]$. The other count pair/mixed intervals also span zero. Wessels count corrections worsen the reference. Raw count and log2 MSE units cannot be compared directly, and neither scale replaces the original endpoint.

Each outer source defines a strong panel by its top 10% mean absolute pair effects. A separate low-single residual-enriched proxy ranks source residual magnitude among genes whose source single magnitudes do not exceed the median. Query pair labels cannot select either mask. All models use the same fold masks; these are diagnostic source-derived panels, not validated emergent gene sets. All nine protocols, all three panels and one/two-held-target subsets are retained.

On Norman count partner folds, the residual proxy gives pair-product and mixed gains of 12.49% and 9.23%, with positive conditional lower bounds. These favorable subsets do not override failed all-gene gates. The top source-selected 10% strong genes contribute 99.17% and 99.43% of Norman count reference error in random and partner scoring, illustrating endpoint concentration.

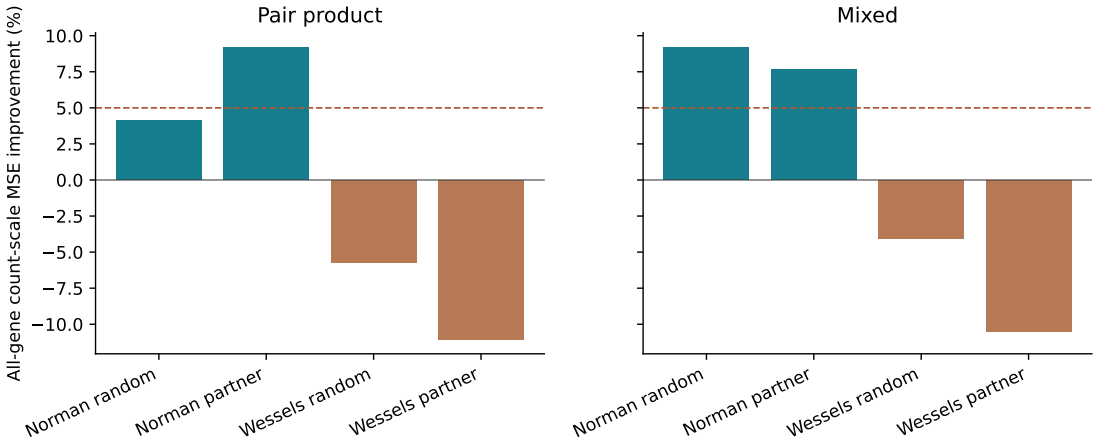


Figure 18: Count-scale point gains on Norman do not pass the paired-uncertainty gate. Wessels remains unfavorable. Full panels, intervals and target subsets are retained.

Identity-only recovery of the original partner assignments separates 120 one-held-target and 28 two-held-target Norman occurrences, versus 58 and 11 on Wessels. These correspond to only 88 and 40 unique conditions: 60 Norman pairs and 29 Wessels pairs recur in two folds. Pairwise intersections and every coverage count are saved; repeated occurrences do not supply independent replication. Opposite technical-half diagnostics retain complete shared reference matrices and report residual MSE, direction cosine and the main/residual/cross-term identity on all panels. Frozen QC, full-reference configuration selection and source-selected masks prevent a biological-independence claim.

Count fitting takes 14.39 elapsed seconds; the complete fitting and diagnostic runner takes 37.89 seconds on one CPU thread, with 243.8 MB completion RSS, not a measured peak. No pair component advances on these primary gates. Independent cultures and sealed state/pair confirmation remain missing.

E.5 Frozen external UPR transfer feasibility

A metadata-qualified local source extends the audit beyond the exposed development banks. Original Adamson UPR counts and full guide barcodes provide matched controls, three exact genetic singles and three pair identities in each of three chemical backgrounds (Adamson et al., 2016). Publisher methods map capture groups to tunicamycin, thapsigargin and DMSO. The ATF6/PERK/IRE1 triple is genuine, but one triple identity supports inventory rather than a general higher-order benchmark. The backgrounds alias capture libraries, and independent culture initiation is unverified.

The fixed observed-single experiment aligns 10,423 genes using source-selected Norman genes and exact query Ensembl metadata. It freezes the Norman library normalizer and source-only three-fold selection over the existing mean, global, summed-single, pair-product and mixed menus. Source-selected global correction and mixed predictions are saved before evaluator access. All nine pair/background outcomes are scored once locally. This is a public-source feasibility pilot, not sealed, powered joint pair/state confirmation.

Model	MSE
additive	0.00184324
mean residual	0.00234538
global correction	0.00292072
single sum	0.00252739
pair product	0.00289898
mixed	0.00264696

The frozen mixed candidate gains 9.37% relative to the source-selected correction, with diagnostic paired MSE-reduction interval $[-0.00015503, 0.00067825]$. Its lower bound includes zero, so the materiality gate fails. Additivity has MSE 0.00184324, lower than the source-selected correction's 0.00292072; a favorable mixed-versus-correction point estimate cannot establish value over this simpler retained competitor. No post-query family replacement or architecture advancement occurs.

Preparation takes 4.95 seconds with 3.08 GB process peak working set. Frozen source fitting takes 1.78 seconds with 370.0 MB completion RSS. Both use one CPU thread and remain below the registered budgets. The crossed target/background bootstrap retains 1773 nonempty draws from 2,000; it holds reference means and selection fixed and supplies no biological confidence interval. CRISPRa-to-CRISPRi, study and chemical/capture shifts remain confounded.

A complete rerun of the unchanged locally prespecified development studies preserves every prior result and records a different numerical environment. The original NumPy2.5.3/SciPy1.18.1/PyTorch2.9.1+cu126 environment differs from the current NumPy2.4.0/SciPy1.16.3/PyTorch2.9.1 CPU environment. Norman neural shared/pair MSEs change from 0.00740358/0.00740706 to 0.00740673/0.00749855, and four inner-selected configurations change. All development gate decisions remain unchanged; classical predictions agree to numerical precision. The audit retains every numerical difference and the failed strict-replay attempt. This is reproducibility evidence under a changed environment, not new independent replication or guaranteed bitwise retraining.

F Separately scoped trainable model-series development

The owner authorized a bounded three-model development cycle on 10 October 2026 after the unfavorable exposed-source Combi-seq benchmark. This amendment does not replace any original numerical result, prospective architecture freeze, biological gate or confirmation requirement. It estimates retrospective aggregate expression-change prediction conditional on measured singles, a declared deposited assay and source stratum. Safe source-specific checkpoints and a common training, inference and evaluation interface are the deliverable; no universal cross-assay weights or unseen-single capability is claimed.

Opal-Base selects and fits the existing thirteen-entry reference subset inside acquired development folds, retaining observed-single additivity and a regularized correction. Opal-Interaction freezes this expert and learns a symmetric residual using a shared two-layer tanh encoder with widths 64 and 16, an elementwise latent product and an output projection. Opal-Guard learns an affine sigmoid mixture from three grouped out-of-fold expert predictions, each with its own nested selection, using graph degree, measured-single distance and source-context support. Unsupported

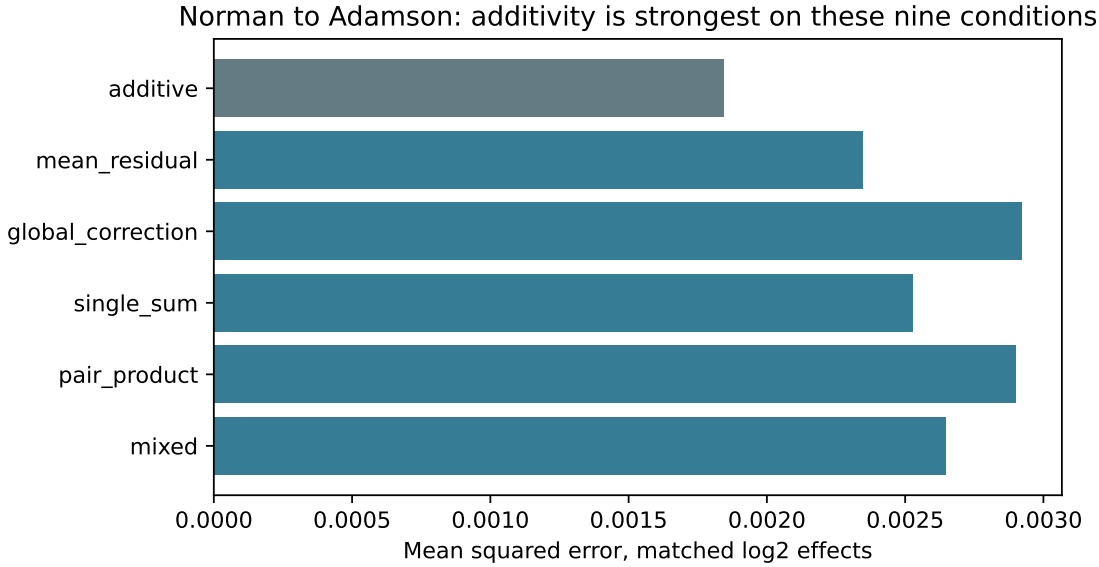


Figure 19: All retained competitors on the same nine external conditions. Additivity is stronger; the frozen mixed candidate fails the paired-uncertainty gate.

components or sample contexts force the declared Base fallback. The mixture is not calibrated uncertainty. Missing assay, ordered genes, units or essential measured-reference inputs causes refusal.

The new configuration fixes one training seed, two penalties, Adam at learning rate 0.003, a 60-epoch branch cap and a 120-epoch gate cap with development-only stopping. The expert and gate budgets differ and are disclosed. All selections, epochs, failures and unavailable results remain saved. Final experts are shared across the variants; every variant has its own inference behavior and safe checkpoint. Held pair outcomes and unacquired labels cannot enter fitting, representation learning or inner validation. Query singles and controls are explicitly allowed.

Method	Combi-seq	PertSpectra TNF-positive	PertSpectra TNF-negative
No change	5.713889	0.137007	0.075095
Additivity	8.993460	0.099904	0.094443
Opal-Base / selected reference	3.872700	0.085929	0.082279
Historical pair-kernel family	4.388534	0.085133	0.082056
Opal-Interaction	3.871518	0.086707	0.083545
Opal-Guard	3.871571	0.086159	0.082651

Table 16: Full-development descriptive MSE for separately fitted source-specific models. Each column uses its own reference-derived 512-gene endpoint and assay contract; absolute losses must not be pooled or compared across columns. Combi-seq equally weights 374 canonical pairs and three reported runs. Each PertSpectra context equally weights its complete pair panel and four source sample labels.

The Combi-seq bank, gene endpoint, canonical outer assignments and historical predictions reproduce exactly. Interaction and Guard reduce MSE from Base by approximately 0.03%, which does not warrant replacing Base. The original pair-kernel result remains unfavorable. Nested uniform and identity-only target-balanced prefixes use 0, 25, 50 and 100% of available development pairs with seeds 11, 29 and 47; inner validation and cross-fitting labels count toward each budget. Outer evaluation populations remain unchanged. At zero labels only no-change and additivity are available. Pair-supervised methods are unavailable when the fixed support rules fail. The intermediate Combi-seq curves favor balanced coverage, while the nonlinear additions change error little. These are retrospective label-coverage curves, not validated active learning or prospective assay savings.

The bounded PertSpectra extension verifies the full 3,142,132,144-byte H5AD by SHA-256 before expression decoding. Its deposited counts layer contains finite nonnegative integer values and is aggregated with bounded CSC column operations, using library-total-10000 normalization over the deposited 13,901 genes and transformation after

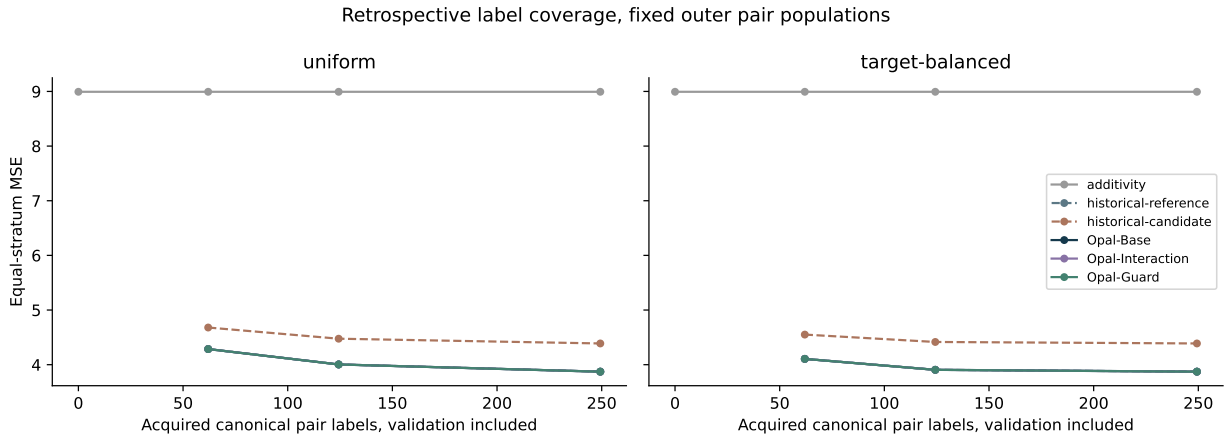


Figure 20: New series label-budget curves generated from saved records. Base and the selected historical reference overlap exactly. Every seed and budget is retained; the curves are descriptive and do not provide biological intervals.

label-conditional means. Controls have one or two control-guide calls; singles include an optional control guide. Ambiguous, intergenic and self-pair assignments are excluded. Gene labels average deposited guide mixtures and do not identify exact matched-guide causal effects. Each source sample/context retains its control, single and pair cell counts and mixtures, requiring at least ten cells and complete reference support across the four sample labels. Source preprocessing remains upstream-dependent. Reconstructing the authors' pooled condition/context count filter from metadata gives exactly the paper's 65,648 cells from 73,015 deposited rows; Opal intentionally uses a separate population and endpoint.

The resulting TNF-positive and TNF-negative panels have 108 and 52 canonical pairs, respectively. Both new nonlinear variants lose to Base in both contexts. No-change also beats every fitted model on the TNF-negative endpoint, so pair supervision is not recommended for that deposited panel. Uniform coverage is more favorable than target balancing at the observed intermediate budgets on these contexts, differing from Combi-seq. All small-budget unavailable cases, three shuffle controls and gate ablations are retained. These findings do not establish biological interaction detection, a noise ceiling, independent-culture confirmation or TNF-state transfer.

A separate leave-SampleIndex-out diagnostic trains on the other three source labels within the same TNF context while excluding every occurrence of each query pair. It uses the held sample's controls and singles, fits transforms and selects configurations within training, and weights training strata equally. Base MSE is 0.100691 in TNF-positive and 0.097187 in TNF-negative, worse than corresponding additivity errors of 0.099904 and 0.094443. Interaction and Guard use their declared Base fallback for unsupported sample contexts and therefore give identical predictions. This is transport across source sample labels, not evidence of independent cultures.

All 432 Combi-seq and 864 within-stratum PertSpectra checkpoints replay saved predictions exactly in the recorded environment. Combi-seq refits reproduce complete selection and epoch histories. Installed-wheel training, batch prediction, evaluator scoring and loading pass outside the checkout. New integrity tests join the existing Windows/Linux Python 3.11/3.14 matrix, but remote CI for these local changes is pending because the owner prohibits commits and pushes. New weights and large arrays stay local, derivative terms remain subject to release review, and the approved private GitHub v0.1.2 assets are unchanged. The independent-culture, empirical biological-power, external registration, correspondence and unapproved publication requirements remain deferred.

G Completed finite model-series evaluation

This section reports the subsequent computational execution separately from the original studies and the preserved first model-series cycle. The public names remain Opal-Base, Opal-Interaction and Opal-Guard. Execution-cycle identifiers describe provenance, not additional named models. All panels were previously exposed development sources. The experiment is retrospective engineering evidence and does not pass the uncompleted independent biological gates.

G.1 Distinct trainable mechanisms and honest selection

Opal-Base selects among 24 main-effect references: the thirteen preserved configurations, no-change, equal-pair response mean, three affine-additivity penalties, and six training-only single-sum PCA ridge configurations. Affine shrinkage constrains the scalar coefficient to $[0, 1.5]$ with penalties $0.1, 1, 10$ toward one. The single-sum PCA ranks are four and eight, with the same penalties. The null is eligible during acquired-label selection; zero acquired pair labels leave all supervised models unavailable.

Opal-Interaction adds explicit symmetric cross-coordinate features. For training-only PCA coordinates u, v, diagonal features are $u_i v_i$ and off-diagonal features are $(u_i v_j + u_j v_i)/\sqrt{2}$. Input ranks four and eight yield ten and 36 coordinates. Training-only orthogonalization removes the selected main-effect design; weighted output SVD uses acquired residuals, with ranks four and eight. Weighted ridge penalties are $0.1, 1, 10$. In-sample and honest cross-fitted residual-target routes give 24 candidates plus a zero branch. Acquired support must supply at least twice the feature dimension, and numerical rank failures remain explicit. Cross-fitted targets come from smaller expert training sets than the final refit; this shift is measured rather than assumed beneficial.

Opal-Guard learns a convex mixture of Base, Interaction and an acquired-pair residual-memory expert. Retrieval standardizes observed singles using training inputs only, compares direct and swapped pair distances, and uses four or eight neighbors. Six support features drive a softmax router. Router penalties $0.1, 1$, fixed seed, 100-epoch cap, ten-epoch patience and fixed stopping tolerance are frozen. Every router-validation identity is excluded globally before constructing any expert OOF row, residual basis, scale or memory bank used in that selection comparison. Expert availability is intersected over nested fitting and final refitting. Fewer than twelve expert-training pairs, four pairs in a partition or 24 router pairs trigger the declared conservative fallback. An expert cannot appear only at final refitting.

The preserved first-cycle router produced expert OOF predictions before choosing its router-validation split. An expert predicting a router-training row could therefore have used another router-validation outcome. This is a verified internal selection dependence, not outer-test leakage. Original predictions and results remain unchanged; the new nested mutation test verifies exclusion before expert construction. Unrecoverable historical gradient diagnostics are recorded as missing.

G.2 Frozen populations, label access and outcomes

The primary populations, scales, ordered 512-gene panels and outer assignments are unchanged from the first cycle. Combi-seq contributes 374 canonical pairs across three reported runs; the separate TNF-positive and TNF-negative contexts contribute 108 and 52 pairs across four SampleIndex strata. These are matching strata, not independently qualified cultures. Nested zero, quarter, half and full prefixes use acquisition seeds 11, 29 and 47 under uniform and identity-only target-balanced orders. Evaluation populations never shrink with acquisition. Two additional fixed training seeds are descriptive full-label checks, not alternatives from which a favorable result is selected.

The primary grid contains 234 Combi-seq and 312 conditions in each TNF context. Coherent historical Norman and Wessels banks add 78 conditions on each log/count scale, totaling 312 retrospective conditions. Their 88 and 40 original development pairs are reused; the protected 38 and 18 original query pairs remain unopened. Full-development historical aggregation draws on 704 Norman or 160 Wessels pair-stratum profiles. Count/log views are two scales of each source, not independent confirmations. Fifty-one pooled, same-context transport and leave-target-group-out sensitivities retain explicit additional label counts; a later completion adds their historical references without acquiring additional pair labels. Fixed 25/50-cell support subsets and alternate normalization-order diagnostics remain separate from the ten-cell primary population. Deposited metadata do not establish guide equivalence or exact physical-dose matching.

All fitting, representation fitting, stopping, OOF experts and router validation use only the declared acquired labels. Label roles overlap and cannot be summed as independent observations. The separate label ledger counts unique pairs and underlying pair-stratum profiles. Outcome-based exact, constant, identity-shuffled and disjoint technical-half controls are evaluator-only. Shared control and single half-draws are reused across affected pairs. Technical halves

describe sampling within this deposit and do not establish a biological noise ceiling, independent-culture uncertainty or power.

Primary scoring remains gene-averaged MSE, equal canonical-pair weighting within each stratum and equal-stratum aggregation. MAE, centered Pearson with undefined zero-norm rows, signed correction magnitude and gene-centered held-profile retrieval are saved separately. Retrieval uses the same-stratum held population and average ranks for ties. Alternate normalization changes the endpoint and is explicitly labeled; it does not replace primary scores. Paired loss distributions and fixed-seed variation are descriptive because targets and reference measurements are shared.

Method	Combi-seq	TNF-positive	TNF-negative
no-change	5.713889	0.137007	0.075095
additivity	8.993460	0.099904	0.094443
historical-reference	3.872700	0.085929	0.082279
Opal-Base	2.111564	0.077789	0.071218
Opal-Interaction	2.109919	0.077723	0.071218
Opal-Guard	2.112194	0.078360	0.071218
Constant mixture	2.107765	0.077688	Unavailable
Matched main effect	2.105071	0.077789	0.071596

Table 17: Full-label primary MSE, uniform acquisition seed 11 and the prespecified training seed. Every method uses the unchanged evaluation population. TNF-negative constant mixing lacks support in four of twelve folds, so its eligible-only mean cannot fill this table. Columns have different units and are not pooled.

G.3 Negative promotion result and scoped retrospective gains

Interaction improves over Base by 0.077882% on Combi-seq and 0.083848% on TNF-positive. Both are below the prespecified 2% practical improvement criterion. Guard loses to Base and constant mixing on these panels. TNF-negative Interaction and Guard use Base fallback. Constant mixing yields MSE 2.107765 and 0.077688 on the two fully supported primary panels; matched main effects yield 2.105071 on Combi-seq, also below Interaction. The larger affine shrinkage and response-mean gains show that improved reference calibration accounts for most primary progress; residual predictions are not causal interaction estimates.

Norman Guard yields MSE 0.006614 versus Base 0.006830 on the log scale, and 1.333030 versus 1.403014 on the count scale. These are retrospective improvements of 3.168524% and 4.988153% on one source. Constant mixing yields 0.006658 and 1.340040. Wessels models coincide at 0.005159 and 0.186866 with declared fallback. No new primary nonlinear default is promoted. Base is the simple fitted primary recommendation; no-change remains available without pair labels. Experimental Norman routing may be inspected locally, but its two scales do not supply two independent positive source panels.

G.4 Executed upstream comparisons and concrete refusals

PertSpectra's interpretable factorization uses label and graph structure (Chang et al., 2026), while TxPert combines multiple knowledge graphs (Wenkel et al., 2026). Strong-baseline evidence (Ahlmann-Eltze et al., 2025) and recent calibrated-metric evidence (Miller et al., 2026) motivate retaining both primary MSE and discrimination diagnostics. Published correlations, outcome-selected DEG weights and technical-duplicate controls cannot populate the matched primary table or enter predictor selection.

Pinned PertSpectra, GEARS and CPA author modules execute 24 full-label TNF folds each. All 72 are protocol adaptations, not full author configurations or native frontier comparators. PertSpectra uses the unchanged factorization core, observed control/single and acquired-pair cells, without the unavailable author graph, under a two-configuration, 100-update search. Mean per-cell transformed inputs differ from Opal's transform of condition means. GEARS uses public GO plus observed-single coexpression, two hidden widths and twenty epochs. CPA uses its native Gaussian aggregate encoder/trainer, two widths, categorical presence coefficients and 120 epochs; aggregate observations are not fabricated independent cells or measured physical doses. GEARS/CPA receive all nineteen observed singles in the stratum in addition to acquired pair labels. Their control/profile input tier differs from pair-extracted observed-single fitting.

Exact licensed PRESAGE GenePT and ESM vectors were already cached and reused. Native probes fail on an undefined PCA decoder, a vector-pooling call signature, and a two-row GAT output for one two-target pair. No invented aggregation or substituted random projection is scored. Public-only TxPert import fails on missing `omegaconf`; its

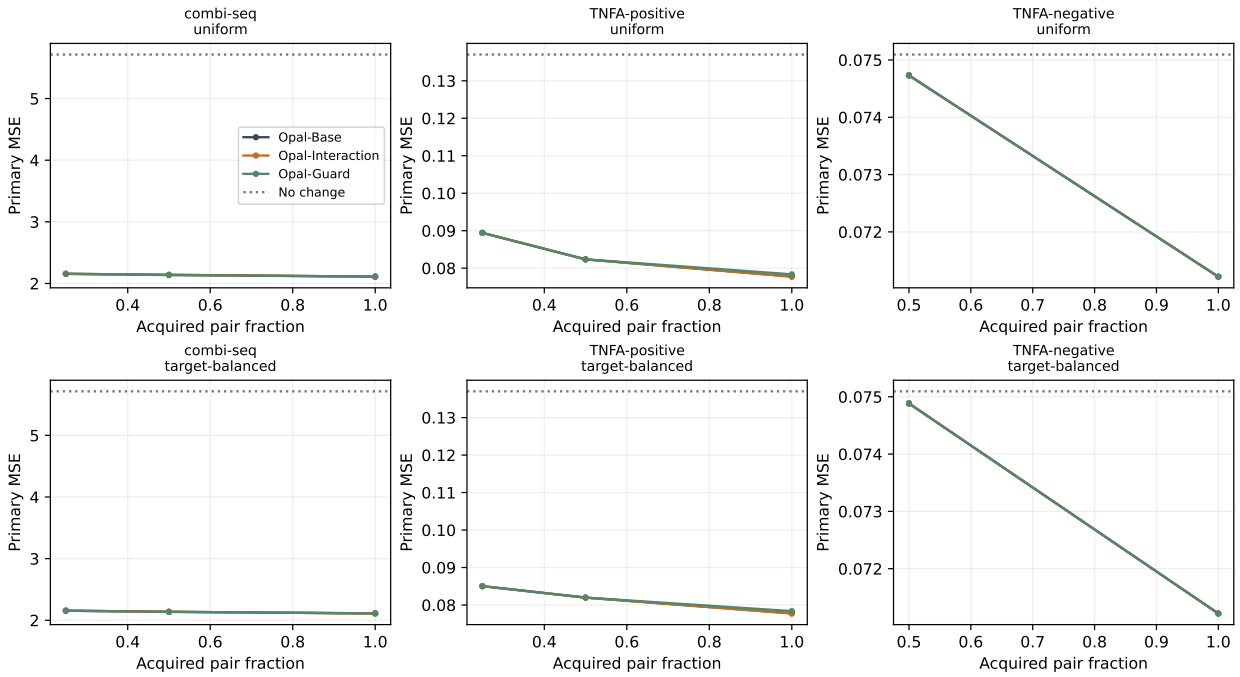


Figure 21: Frozen primary acquisition curves, averaging the three acquisition seeds within each order. Missing zero-budget supervised predictions stay unavailable. The same outer pair population is scored at every prefix; no-change is shown explicitly. Complete retrospective curves and all unavailable low-support cases are retained in compact saved records.

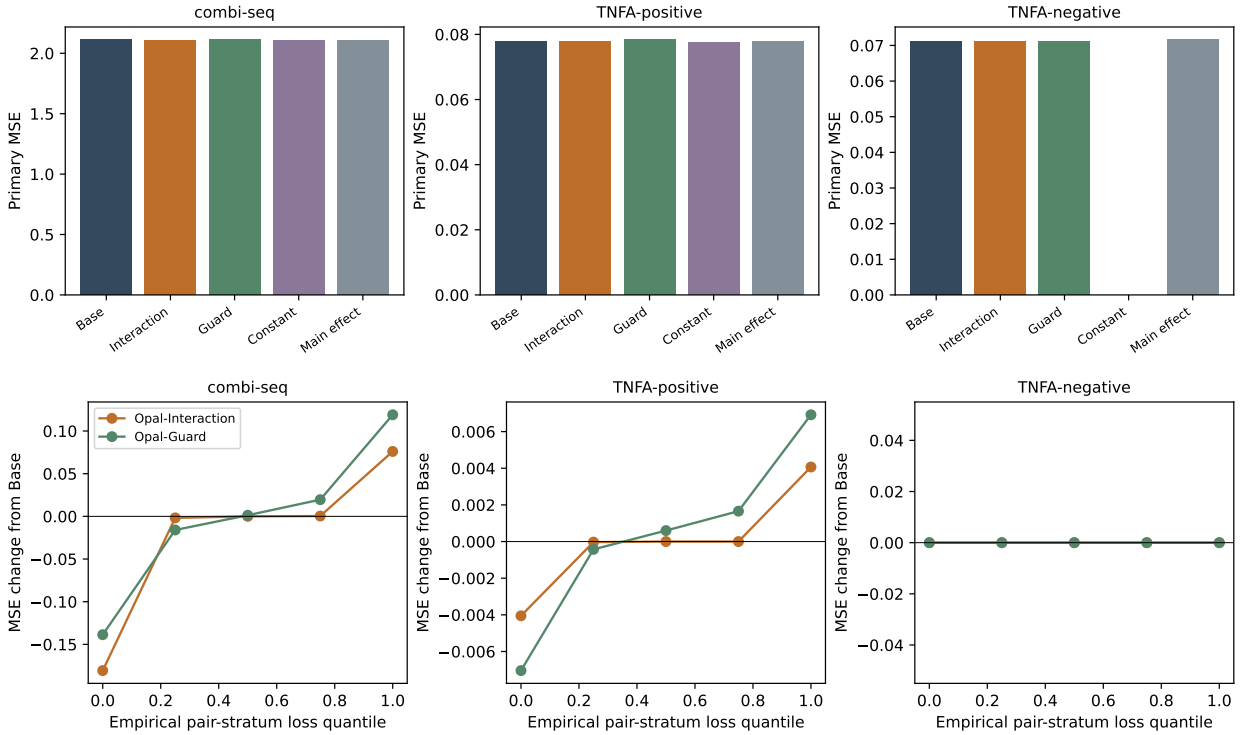


Figure 22: Primary full-label method comparison and empirical five-number summaries of paired changes from Base. The missing TNF-negative constant bar means incomplete expert availability. Pair-stratum loss distributions are not independent-pair confidence intervals.

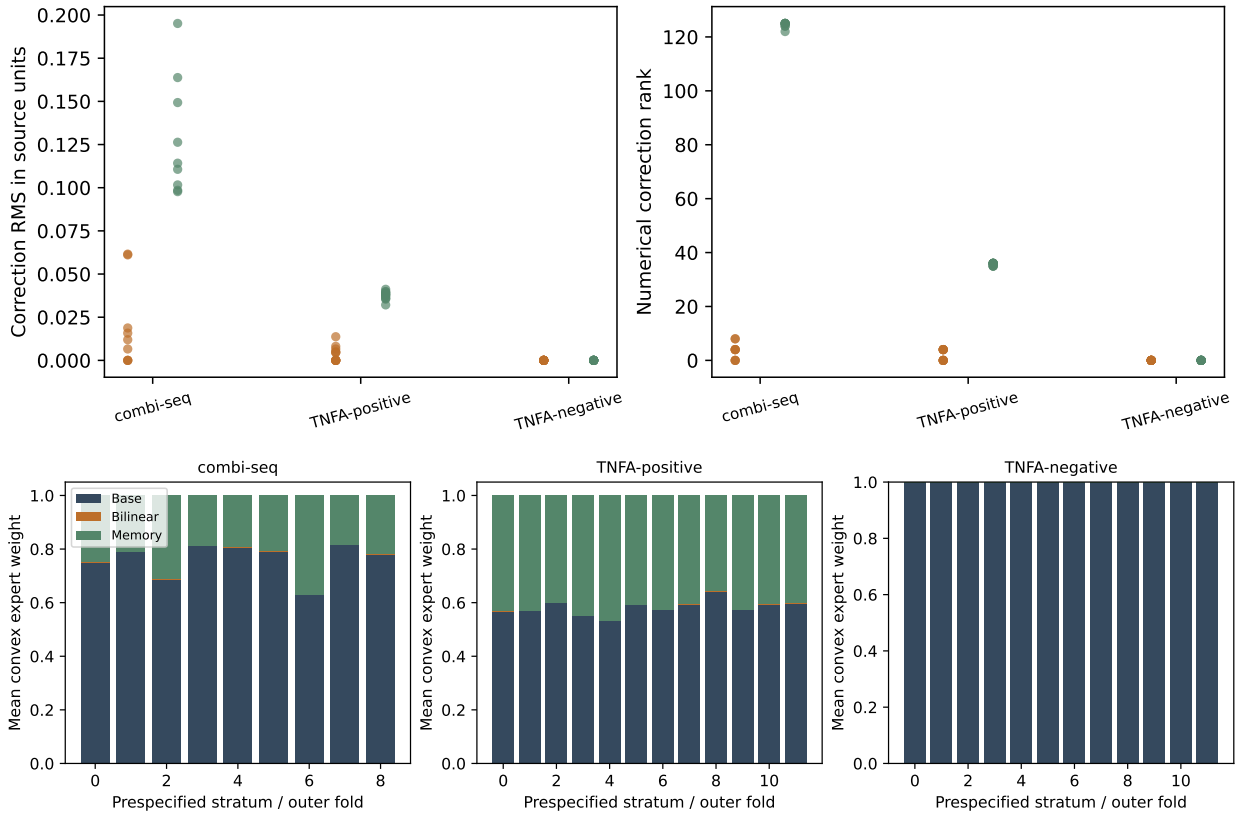


Figure 23: Saved-checkpoint correction magnitude/rank and mean convex expert weights. Zero corrections and unavailable experts remain visible. A nonzero branch or memory weight demonstrates a fitted mechanism, not biological interaction recovery.

Executed adaptation/reference	TNF-positive MSE	TNF-negative MSE
PertSpectra	0.182326	0.126928
GEARS	0.093978	0.079552
CPA	0.113226	0.071731
Linear, same extra observed singles	0.082463	0.070020
Opal-Base, primary information tier	0.077789	0.071218

Table 18: Actually executed comparisons on unchanged primary endpoints. The extra-single linear control uses the same nineteen singles, one control and acquired pairs as the aggregate adaptations, with three development-selected penalties. It improves on Base in TNF-negative. Input tiers and search budgets differ, so these numbers do not rank full published configurations.

supported Python/Torch closure is unresolved in the existing environments, and proprietary graphs are unavailable. Exact attempts and source pins are retained. Missing canonical inputs and unsupported closures prevent corresponding native competitiveness claims. State/CellFlow remain audited distributional context without copied published scores.

G.5 Cost, usable checkpoints and limits

The search is unequal: 24 Base references, 24 Interaction configurations plus zero, four retrieval/router settings with fully nested refits, and two native-adaptation configurations followed by selected refitting. CPU fitting uses one thread and float64; upstream adaptations use the existing Python 3.10/Torch 2.0 CUDA environment on an RTX 4060. Train/inference time, checkpoint bytes and fitted coefficient, representation, memory and support values are saved separately. Native training RAM was not captured; replay RAM is labeled reconstruction/inference RAM and cannot complete the efficiency criterion. CPU/GPU timing is not an equivalent-hardware Pareto comparison.

Safe JSON/NPZ evaluation checkpoints and content-addressed numeric-row stores retain source, scale, ordered genes, reference provenance, split, selected penalties, support and expert attribution. All evaluation checkpoints replay; the 72 upstream adaptations replay within the frozen 10^{-6} float32 tolerance, while Opal float64 replay uses 10^{-10}. Historical 1,296 budget and 72 transport checkpoints also replay exactly. Later all-development refits have separate deployment manifests and no inherited evaluation score. Batch inference is allowed for deployment checkpoints; the sealed retrospective evaluator refuses them. Weights remain local because source/publication authorization for redistribution is absent. The local review package is 0.1.3.dev0; approved 0.1.2 assets are unchanged.

The owner-approved retained-artifact ceiling is 3 GB, with a 1 GiB reserve on writable C. D/E are read-only below reserve; no old files are removed. The initial 1 GB and later 2 GB pauses, retries and accounting correction remain recorded. Default promotion fails, and complete native competitiveness and efficiency criteria remain unestablished. Independent cultures, matched reference/dose provenance, biological precision/power, custody and actual laboratory validation remain external dependencies.

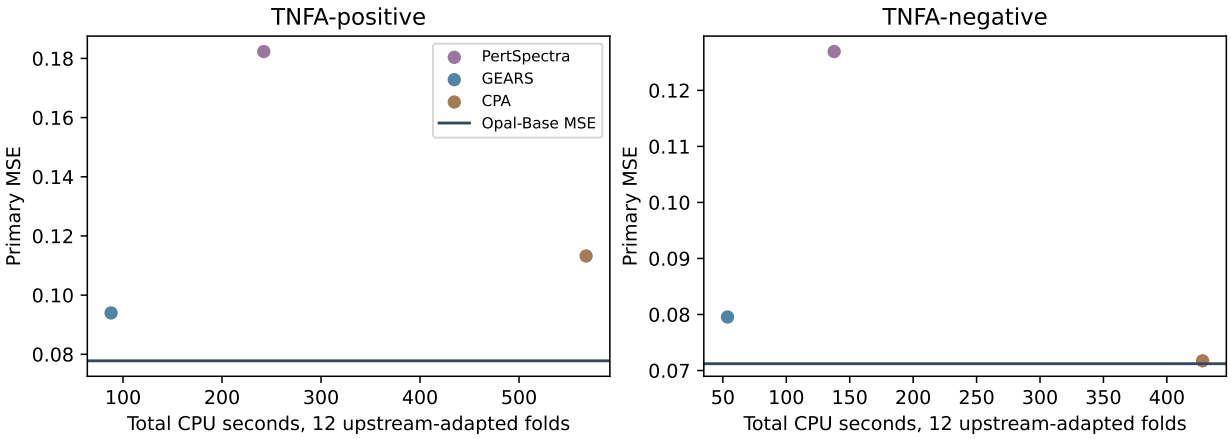


Figure 24: Actually executed upstream adaptations, total process CPU seconds across twelve folds in each TNF context versus primary MSE. The horizontal line is Base's MSE, not a timing comparison. GPU work and unavailable native training RAM prevent equivalent-hardware efficiency or a complete Pareto claim.

H Primary diagnostic figures

The following plots are generated from frozen primary artifacts. Capacity compares registered candidates rather than an exhaustive capacity search. Learning curves keep full-development-selected hyperparameters fixed; split-half spectra describe technical measurement reproducibility, not biological replication or calibrated risk.

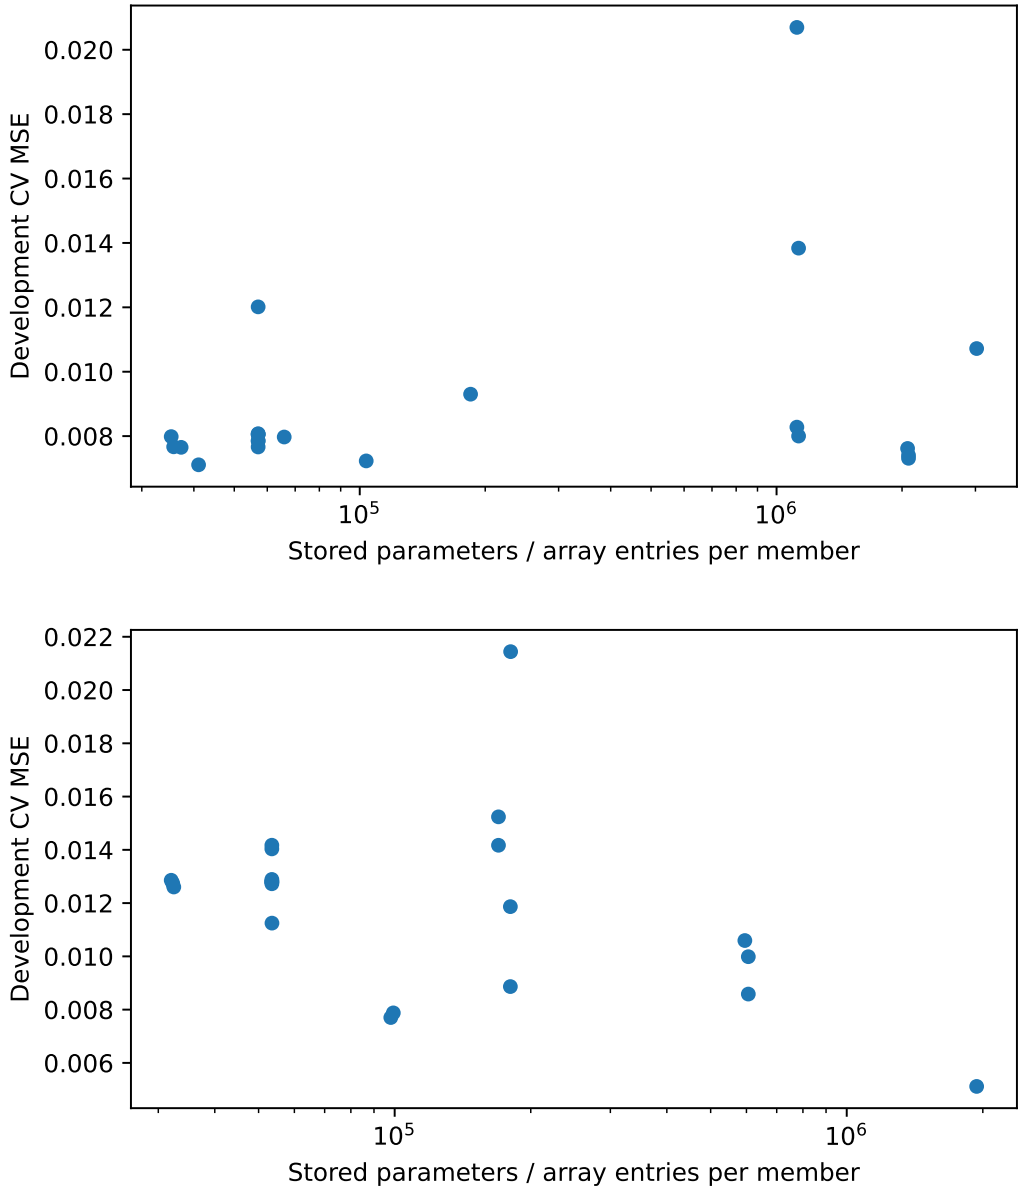


Figure 25: Registered capacity comparisons for Norman (top) and Wessels (bottom). Parameter counts do not include complete reference-bank and library memory.

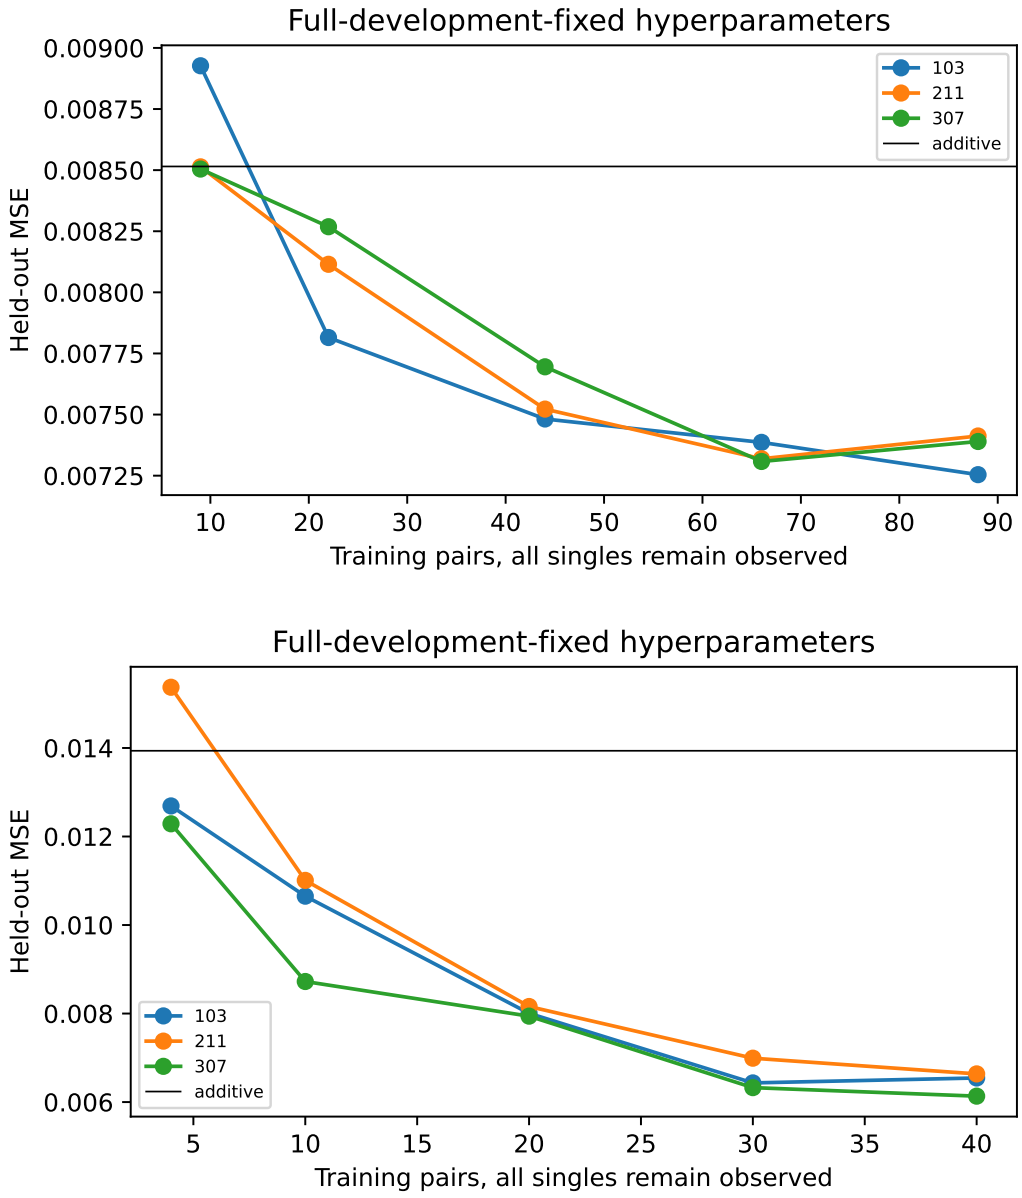


Figure 26: Fixed-hyperparameter development-subset sensitivity for Norman (top) and Wessels (bottom). Three nested subset orders share the same held-out evaluation and are correlated.

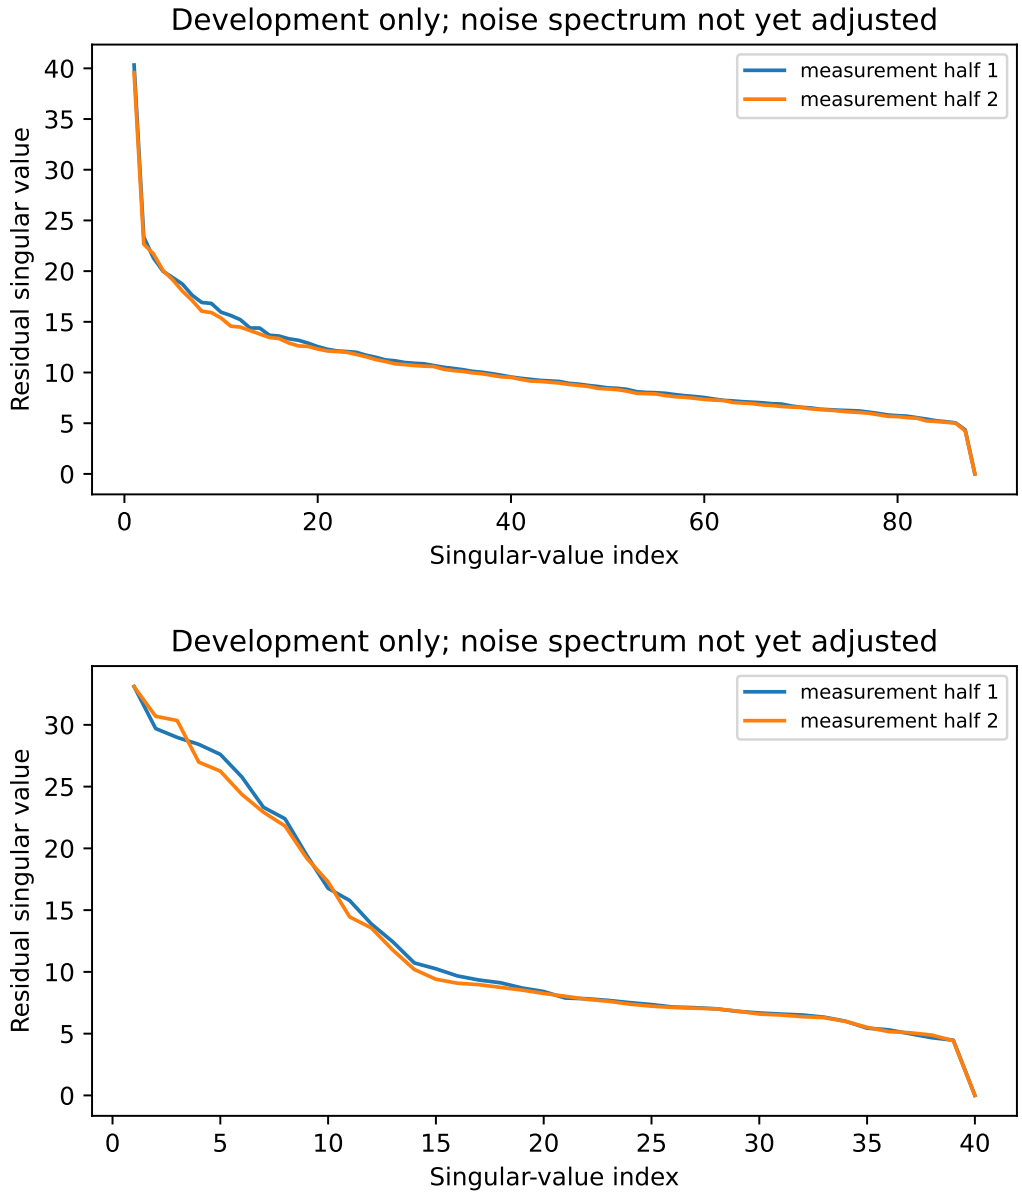


Figure 27: Development residual spectra from the two technical halves for Norman (top) and Wessels (bottom). Noise has not been subtracted; these do not identify biological interaction mechanisms.

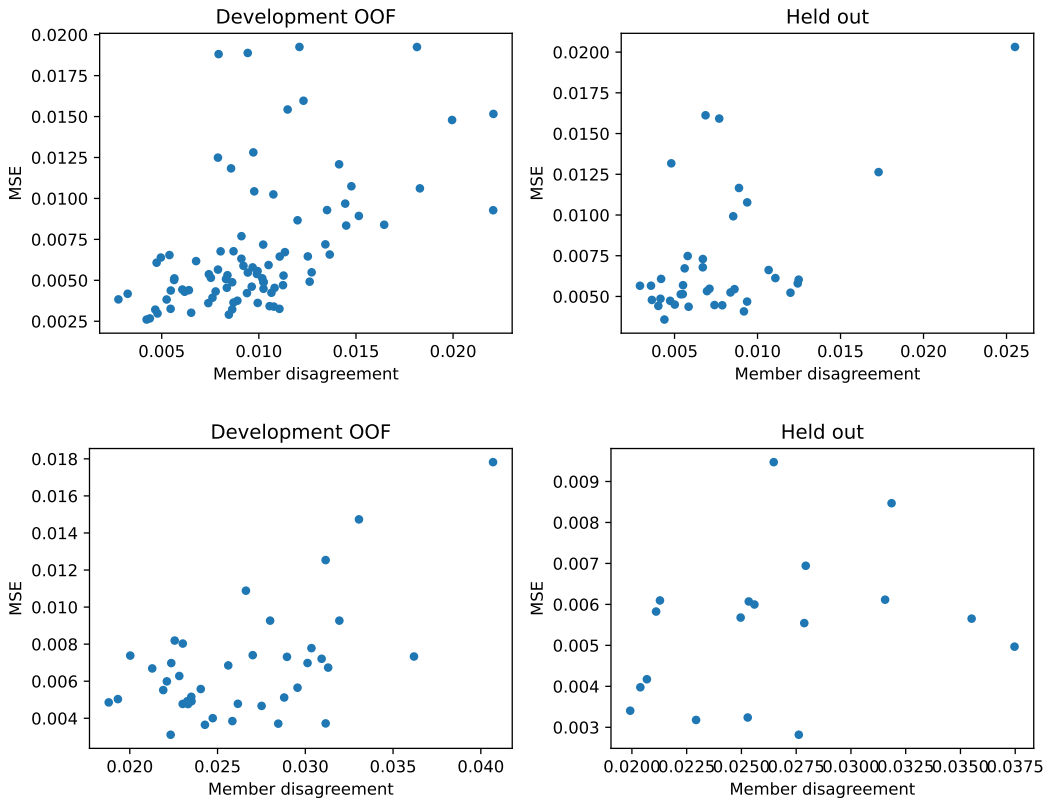


Figure 28: Uncalibrated ensemble-spread versus condition-error diagnostics for Norman (top) and Wessels (bottom), shown on development out-of-fold and held-out predictions. Association does not establish a calibrated failure probability.

I Measured-single variance information: bounded development study

A subsequent frozen experiment asks whether projected variance of measured single-cell singles adds held-partner mean-response information. Both exposed PertSpectra contexts retain their original ordered 512 genes and log2 library-10000 aggregate endpoint. Control-cell-only projections, training-only output bases and grouped nested ridge selection compare means, capacity-matched squared means, real variance and compatible-stratum shuffled variance, each with twelve rank/penalty configurations plus zero correction. A shared ten-cell reference draw is a separate precision sensitivity. No pair cells enter information-model features; query singles remain permitted, so partner holdout is not unseen-intervention prediction.

Context	Split	Base	Mean	Capacity	Variance	Shuffle
TNF-positive	Pair	0.07822245	0.07738171	0.07736861	0.07745757	0.07762722
TNF-positive	Partner	0.08275970	0.08254379	0.08253932	0.08256839	0.08260062
TNF-negative	Pair	0.07146807	0.07101187	0.07097175	0.07104495	0.07111463
TNF-negative	Partner	0.07362159	0.07362159	0.07362159	0.07362159	0.07362159

Table 19: Exposed development only, equal-pair within stratum and equal-stratum MSE. Original populations are retained; technical sample labels are not independent cultures. TNF-negative partner heads all use the declared Base fallback because nested selection lacks support.

The primary TNF-positive partner variance gain over Base is 0.231163%, while both equally tuned mean controls perform better. Ten-cell partner variance gains are 0.451677% and zero in the separate contexts. The prespecified administrative gate requires at least 2% improvement over Base and every matched control in both contexts and the precision sensitivity; it fails. Base is retained, without additional models, seeds, endpoints or sources. TNF-negative partner information is inconclusive because support is insufficient, not because variance information is absent. No fresh-source or independent-culture confirmation was executed.

The study retains 96 fold records, 160 canonical pairs and 640 pair-stratum profiles, with 2,560 repeated evaluation events. It summarizes 23,883 control/single cells, not independent cultures. All evaluation heads replay within 10^{-10} and the existing evaluator rejects changed inputs, units, deployment roles and predictions. The full local Python 3.14 suite passes 282 tests. Fresh training and separate replay profiles are available for Opal and three upstream adaptations; historical native training RAM remains unknown. Zero fully eligible native comparisons are established, and CPU/GPU/input-tier differences preclude an equivalent-workload Pareto claim. Complete CPU certification remains unavailable because bootstrap/tool work and polling losses are not measured. Qualified cultures, dose/reference lineage, custody, empirical biological power and exact-revision remote CI remain dependencies.

The previous completed model-series evidence and negative results are unchanged. The new information study does not establish biological interaction recovery, calibrated uncertainty, context transfer, prospective acquisition savings or frontier superiority. Compact provenance and every arm/failure are preserved in `results/opal_information_v1`; detailed eligibility and resource scopes accompany its local report.

J Fixed-context successor implementation and source refusal

Following the preserved negative and source-incompatible bridge evidence, the research default changes to a separately named fixed-context diagnostic, `opal_context_contrast_v1`. Its software is implemented; an owner-authorized real-source preparation attempt on 14 October 2026 refused at exact-assignment qualification before expression access. Scientific fitting and biological response scoring have not been executed. The frozen bridge configuration, its 207-common-selected-gene refusal against the unchanged 256-gene minimum, synthetic failures and original evidence remain intact. They do not establish biological failure of composition transfer.

Both observed TNF contexts supply training information. Globally held canonical pair identities are excluded across contexts. A new shared panel of at most 512 literal genes is selected from equal-context within-context measured-single-effect variance, requiring at least 256 eligible genes and 36 complete pairs. Full-library normalization precedes restriction, condition means precede log transformation, and exactly matched controls are subtracted. Unknown reagent, dose, time, rights or reference support cannot be inferred from source titles.

For permitted single effects m_{ic}, the four-coordinate representation is $z_{ic} = P^T m_{ic}$, with $d_i = z_{i1} - z_{i0}$ and $\bar{z}_i = (z_{i1} + z_{i0})/2$. The contrast target is $t_p = r_{p1} - r_{p0}$, where $r_{pc} = \Delta_{pc} - m_{Ac} - m_{Bc}$. The candidate features are

$$\phi_8(p) = [d_A + d_B,\ |d_A - d_B|], \tag{9}$$

$$\phi_{18}(p) = [\phi_8(p),\ \mathrm{vech}(\mathrm{sym}(\bar{z}_A d_B^T + \bar{z}_B d_A^T))]. \tag{10}$$

Off-diagonal coordinates use the fixed $\sqrt{2}$ convention. Training-only nuisance cross-fitting precedes the rank-two output basis, normalized by $B^T B/G = I$. The 36-coefficient head minimizes mean pair/gene contrast error plus $\lambda\|W\|_F^2$, with the fixed menu $\{0.1, 1, 10\}$ selected by full-response inner validation. A pooled residual kernel k and shared contrast h reconstruct responses as

$$\widehat{\Delta}_{p0} = m_{A0} + m_{B0} + k(p) - (h + BW\phi_{18}(p))/2, \qquad \widehat{\Delta}_{p1} = m_{A1} + m_{B1} + k(p) + (h + BW\phi_{18}(p))/2. \tag{11}$$

All encoders, scalers, bases, nuisance estimates and references are refitted inside three globally grouped outer and inner folds. Required references include fold-refitted Base, calibration/incidence models, phi8, matched-capacity mean/context heads and pooled, per-context and joint-context kernels. Alignment controls retain identical nuisance freedom. Unsupported fits retain complete-population locked-reference fallback and cannot advance.

The evaluator is separate from the sealed prediction stage. It implements context-aggregate response MSE, secondary diagnostics, 99 evaluator-only retrieval-null permutations and 2,000 conditional intervention-product draws. Two disjoint technical-reference role swaps retain the same query identities and pair estimator. Technical separation does not create independent cultures or biological uncertainty. No program utility is claimed without reviewed frozen weights; independent confirmation requires actual culture/assignment/reference evidence, an empirically reviewed useful margin and power, and a separate held-context registration.

The software inventory includes all nested fits, inference banks, failed attempts and replay assets rather than counting the head alone. The proposed small retained budget has not been shown sufficient, and a complete numerical layout and owner authorization are now recorded, but source eligibility remains required. Fixture checks establish algebra, role isolation, refusal, fallback and restoration behavior only. This appendix adds no successor score, biological confirmation, native-method superiority or frontier claim; all historical evidence tables remain unchanged.

The cache metadata supplies 13,901 literal genes and 51 nominally covered pairs, but zero pairs have certified exact assignment equivalence. Inspection of the author-posted data license and supplemental methods resolved local research rights without supplying dose, duration, independent initiation or validated guide-equivalence records. The preparation process stopped as source-ineligible, using 2.84375 CPU seconds and 98,516,992 peak resident bytes. These are qualification costs, not fitting benchmarks or biological measurements. All five advancement gates remain false.

The access-dependent single-response pivot received an outcome-blind, revision-pinned Tahoe catalog audit: 1,344 sample records were inspected using 369,683 transferred metadata bytes, without expression access. Catalog contexts do not certify measured crossed edges; duration, initiation, cell support and a bounded expression subset remain unqualified. A nonquery crossed-repeatability pilot specification is drafted, but no partner measurements, useful margin, power or costs are established. The separate confirmation readiness check reports measurement-not-ready across all 14 evidence gates. These administrative and source-qualification results add no predictor score or confirmation claim.

K Executed crossed-single feasibility alternative

Following the pair-source refusal, a separately registered exposed single-response study predicts missing compound effects in an observed target cell line from measured effects in the other two lines and target training-identity calibration. It uses 96 lexicographically selected compounds, 1,000 nM, 24 hours and a 512-gene control-only panel. Per-cell log normalization preserves the declared single-response endpoint, distinct from the pair diagnostic. Three grouped outer and inner folds retain all nested fits; a fixed 16-coordinate ridge is compared with calibrated mean/difference profiles, a residual kernel and a matched shuffled-feature control.

Target	Candidate MSE	Selector MSE	Difference ridge MSE	Gain (%)
A549	0.03816697	0.03918590	0.03825730	2.600
K562	0.03316807	0.03556390	0.03478723	6.737
MCF7	0.04300216	0.04404643	0.03825636	2.371

All 684 logical invocations, 288 compound-by-target predictions and exact saved-state replays completed. The development spending filter failed. Whole-compound paired sensitivity is conditional on exposed measurements; it is not independent culture or refitting uncertainty. No pair-composition, unseen-context, biological-confirmation or frontier claim follows. Actual compute was 520.766 CPU seconds, with 91,827,750 retained bytes at the receipt. The registered next action is close finite rank-16 probe; no rescue sweep.

L Official chemical measurement qualification

The official GSM7056151 count, annotation and hash files were acquired for a separate technical measurement-reliability pilot. The author methods and RT allocation supply physical cell-origin and 72-hour, matched-dose treatment records; they do not certify independent culture initiation. At the registered 1 micromolar dose, the trametinib star contains 22 canonical pairs across three origins, below the unchanged 36-pair diagnostic floor. All 66 declared pair/context rows are accounted for: 47 supported and 19 unsupported.

Origin	Supported	Response MSE	Zero response	Residual MSE	Zero residual
A172	18/22	0.084340	0.632093	0.170531	0.169099
T98G	19/22	0.069603	0.295970	0.171433	0.125941
U87MG	10/22	0.419838	1.116216	1.015662	0.594808

These averages describe only supported technical-cell partitions. Response agreement improves on zero, but residual-half error exceeds the zero-residual baseline in each origin. No predictor was fit, no biological failure is established, and no confirmation or frontier claim follows. Exact source matching, independent reference/culture precision and empirical biological review remain necessary. Acquisition transferred 611,638,699 bytes; the completed pilot used 974.156 CPU seconds. Earlier preparation failures and incomplete CPU accounting for a controlled stop are preserved separately.

M Retained-output evidence audit

A separate local audit on 15 October 2026 verifies retained assets without fitting or reopening historical protected evaluators. For the exposed crossed-single study, all 594 inner validation losses, selected penalties and reference families reproduce; all 90 final saved-state predictions replay exactly. Recalculated full-population MSEs and the registered conditional sensitivity bounds agree with the reports. The spending filter remains failed. Logical prediction-seal ordering and content hashes do not provide an independent timestamp attestation of historical outcome access or biological replication.

The GxE source hashes and all 66 reported row records verify, and supported-subset summaries reproduce. Processed single/control/pair aggregates were not retained, so inexpensive independent raw-to-metric replay remains incomplete. Shared references, jointly selected reference panel and cell parity couple these technical comparisons; they cannot estimate independently initiated culture covariance. Complete hash-field semantics and dataset redistribution rights remain unresolved. Sampled RSS is not a continuous process peak, and the interrupted preparation attempt retains incomplete final CPU accounting.

Historical bridge custody checks find all 5,142 inventoried files present and hash-matched. Numerical-source matching across the manuscript is a traceability aid, not a certification of comparator semantics, all historical procedures or independent biological claims. The next scientific objective is a separately reviewed, nonquery independent-culture/reference precision pilot with actual lineage, useful margin, power and resource evidence. No predictor is nominated, no physical confirmation is completed and no frontier or mechanism claim follows from this audit.